

Machine Failure Prediction - Data Analytics Capstone Project

This notebook walks through a full machine learning pipeline to predict machine failures using sensor data.

Data Overview

Project Name: *Data Analytics Major Project – Predictive Maintenance using Sensor Data*

Objective: To predict machine failures in advance using historical sensor readings from various machines, enabling proactive maintenance and minimizing downtime.

Dataset Summary: The dataset consists of multiple sensor readings, environmental parameters, and operational metrics from machines, along with a binary indicator of whether a failure occurred.

Number of Features: 11 (including the target variable) **Target Variable:** `Fail` (1 = Failure, 0 = No failure)

Feature Descriptions

| Feature | Description |
|--------------------|---|
| Footfall | Number of people or objects passing near the machine. |
| Temp Mode | Temperature mode/setting of the machine. |
| AQ | Air Quality Index near the machine. |
| USS | Ultrasonic sensor readings (proximity measurements). |
| CS | Electrical current usage (current sensor readings). |
| VOC | Volatile organic compounds level near the machine. |
| RP | Rotational position or RPM (revolutions per minute) of machine parts. |
| IP | Input pressure to the machine. |
| Temperature | Actual operating temperature of the machine. |
| Fail | Binary indicator of machine failure (1 = Failure, 0 = No failure). |

Potential Use Cases

- **Predictive Maintenance:** Forecast failures before they occur to schedule timely servicing.
- **Operational Efficiency:** Optimize operating conditions to extend machine life.
- **Safety Monitoring:** Detect unsafe conditions based on environmental and sensor readings.
- **Cost Reduction:** Minimize unexpected breakdowns and repair costs.

Let's Begin

Import Libraries

```
# Step 1: Import Libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix
import warnings
warnings.filterwarnings('ignore')
```

Load Dataset

```
# Step 2: Load Dataset
df = pd.read_csv('/content/data (1).csv')
```

Dataset First View

```
# Dataset First Look
df.head()

{"summary": {"name": "df", "rows": 944, "fields": [
    {"column": "footfall", "properties": {"dtype": "number", "std": 1082, "min": 0, "max": 7300, "num_unique_values": 99, "samples": [370, 170, 88], "semantic_type": "\\", "description": "\n", "tempMode": "\n", "properties": {"dtype": "number", "std": 2, "min": 0, "max": 7, "num_unique_values": 8, "samples": [1, 5, 7], "semantic_type": "\\", "description": "\n\n"}, "column": "AQ", "properties": {"dtype": "number", "std": 1, "min": 1, "max": 7, "num_unique_values": 7, "samples": [7, 3, 6], "semantic_type": "\\", "description": "\n\n"}, "column": "USS", "properties": {"dtype": "number", "std": 1, "min": 1, "max": 7, "num_unique_values": 7, "samples": [7, 3, 6], "semantic_type": "\\", "description": "\n\n"}]}
```

```

[\"n      1,\n      3,\n      7\n    ],\n  \"semantic_type\": \"/\",\\n  \"description\": \"/\\n      \",\n  },\\n  {\n    \"column\": \"CS\",\\n    \"properties\": {\n      \"dtype\": \"number\",\\n      \"std\": 1,\n      \"min\": 1,\n      \"max\": 7,\n      \"num_unique_values\": 7,\n      \"samples\": [\n        6,\n        5,\n        2\n      ],\n      \"semantic_type\": \"/\",\\n      \"description\": \"/\\n      \",\n      \"column\": \"VOC\",\\n      \"properties\": {\n        \"dtype\": \"number\",\\n        \"std\": 2,\n        \"min\": 0,\n        \"max\": 6,\n        \"num_unique_values\": 7,\n        \"samples\": [\n          6,\n          1,\n          5\n        ],\n        \"semantic_type\": \"/\",\\n        \"description\": \"/\\n      \",\n        \"column\": \"RP\",\\n        \"properties\": {\n          \"dtype\": \"number\",\\n          \"std\": 16,\n          \"min\": 19,\n          \"max\": 91,\n          \"num_unique_values\": 71,\n          \"samples\": [\n            34,\n            36,\n            53\n          ],\n          \"semantic_type\": \"/\",\\n          \"description\": \"/\\n      \",\n          \"column\": \"IP\",\\n          \"properties\": {\n            \"dtype\": \"number\",\\n            \"std\": 1,\n            \"min\": 1,\n            \"max\": 7,\n            \"num_unique_values\": 7,\n            \"samples\": [\n              3,\n              4,\n              1\n            ],\n            \"semantic_type\": \"/\",\\n            \"description\": \"/\\n      \",\n            \"column\": \"Temperature\",\\n            \"properties\": {\n              \"dtype\": \"number\",\\n              \"std\": 5,\n              \"min\": 1,\n              \"max\": 24,\n              \"num_unique_values\": 24,\n              \"samples\": [\n                9,\n                17,\n                1\n              ],\n              \"semantic_type\": \"/\",\\n              \"description\": \"/\\n      \",\n              \"column\": \"fail\",\\n              \"properties\": {\n                \"dtype\": \"number\",\\n                \"std\": 0,\n                \"min\": 0,\n                \"max\": 1,\n                \"num_unique_values\": 2,\n                \"samples\": [\n                  0,\n                  1\n                ],\n                \"semantic_type\": \"/\",\\n                \"description\": \"/\\n      \",\n                \"column\": \"\",\\n                \"properties\": {\n                  \"dtype\": \"dataframe\",\\n                  \"variable_name\": \"df\"\n                }\n              }\n            }\n          }\n        }\n      }\n    }\n  }\n}

```

Dataset Rows & Columns count

```
# Dataset Rows & Columns

df.shape

(944, 10)
```

Dataset Information

```
# Dataset Info

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 944 entries, 0 to 943
Data columns (total 10 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   footfall    944 non-null    int64  
 1   tempMode    944 non-null    int64  
 2   AQ          944 non-null    int64  
 3   USS         944 non-null    int64  
 4   CS          944 non-null    int64  
 5   VOC         944 non-null    int64  
 6   RP          944 non-null    int64  
 7   IP          944 non-null    int64  
 8   Temperature 944 non-null    int64  
 9   fail        944 non-null    int64  
dtypes: int64(10)
memory usage: 73.9 KB
```

Duplicate Values

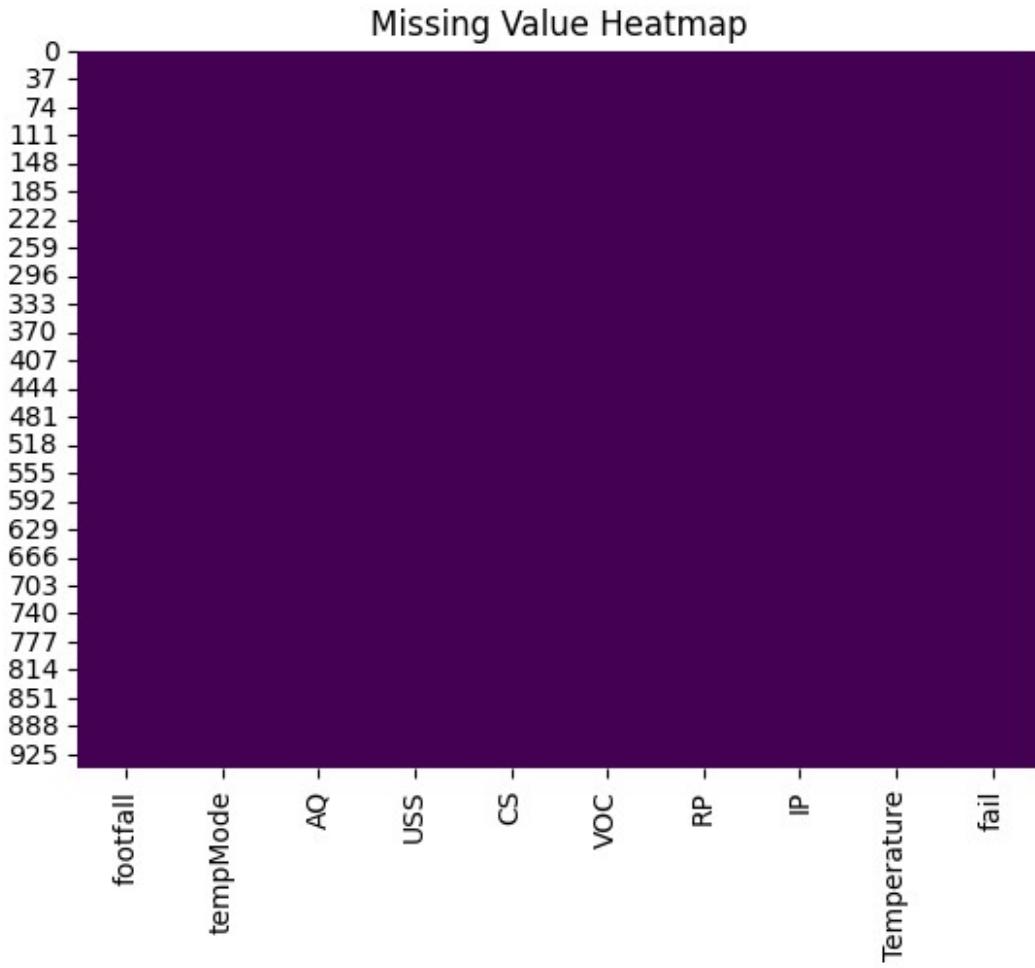
```
# Dataset Duplicate Value Count
df.duplicated().sum()
np.int64(1)
```

Missing Values/Null Values

```
# Missing Values/Null Values Count
df.isnull().sum()

footfall      0
tempMode      0
AQ            0
USS           0
CS            0
VOC           0
RP            0
IP            0
Temperature   0
fail          0
dtype: int64

# Visualizing the missing values
sns.heatmap(df.isnull(), cbar=False, cmap='viridis',
            xticklabels=df.columns)
plt.title('Missing Value Heatmap')
plt.show()
```



Data Cleaning

```
# Step 3: Data Cleaning
print(df.info())
print(df.isnull().sum())
# Fill missing numeric values with median
df.fillna(df.median(), inplace=True)
```

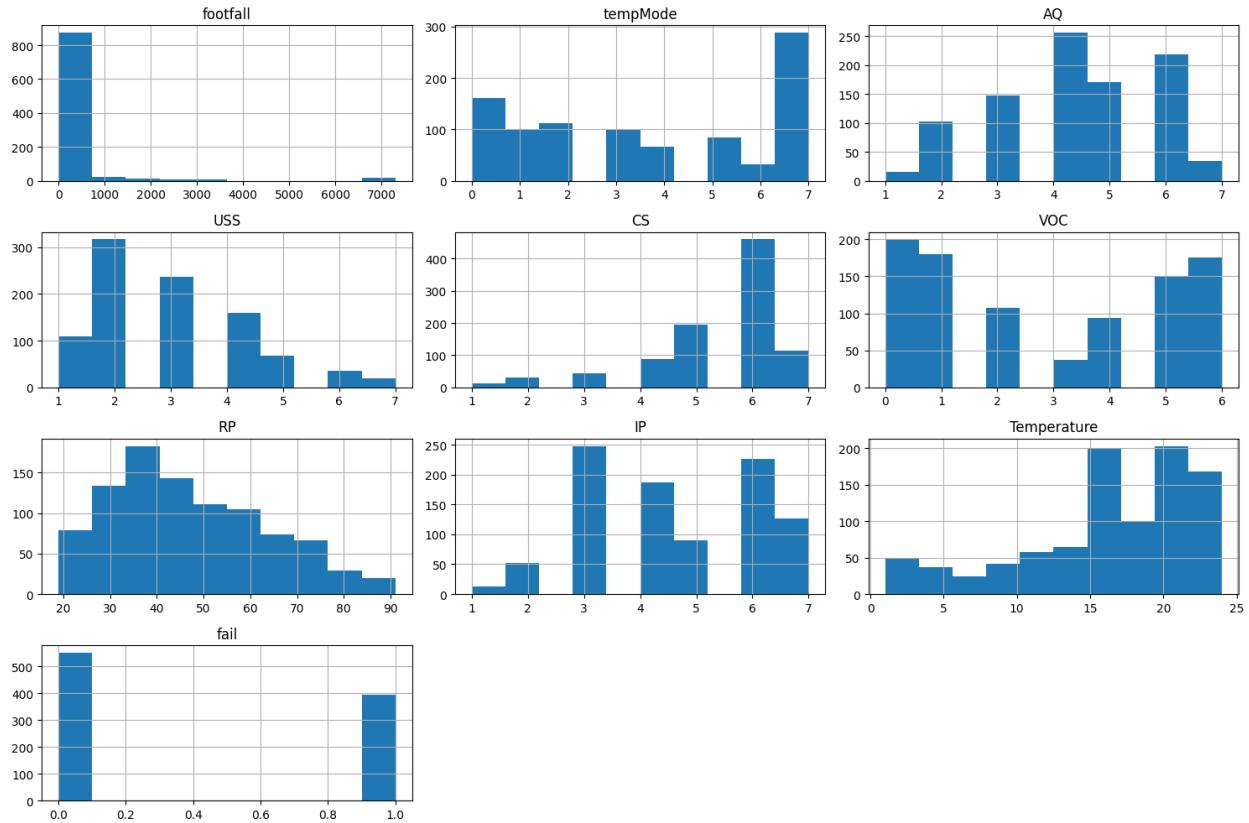
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 944 entries, 0 to 943
Data columns (total 10 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   footfall    944 non-null    int64  
 1   tempMode    944 non-null    int64  
 2   AQ          944 non-null    int64  
 3   USS         944 non-null    int64  
 4   CS          944 non-null    int64  
 5   VOC         944 non-null    int64
```

```
6    RP          944 non-null    int64
7    IP          944 non-null    int64
8  Temperature  944 non-null    int64
9    fail        944 non-null    int64
dtypes: int64(10)
memory usage: 73.9 KB
None
footfall      0
tempMode       0
AQ             0
USS            0
CS             0
VOC            0
RP             0
IP             0
Temperature    0
fail           0
dtype: int64
```

Data Visualisation

Distribution of each numerical feature

```
# Visualize the distribution of each numerical feature
df.hist(figsize=(15, 10))
plt.tight_layout()
plt.show()
```

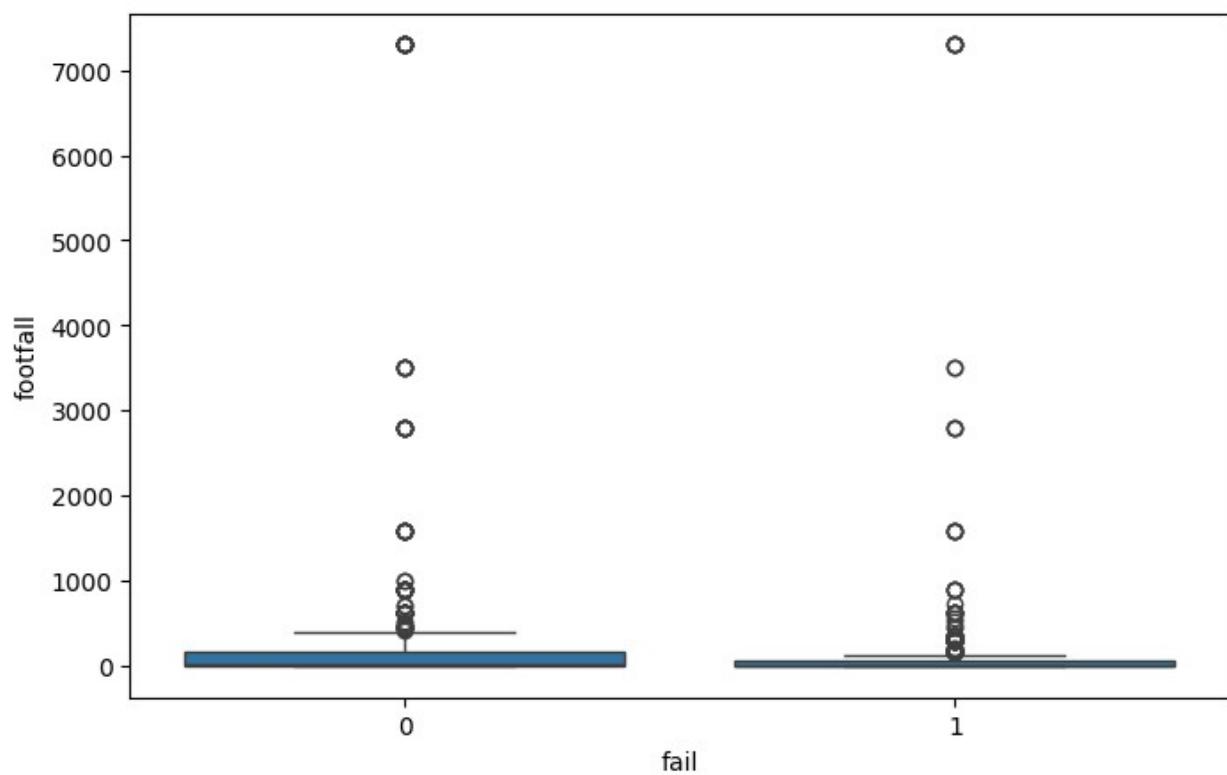


Relationship between features and failure (Box plots)

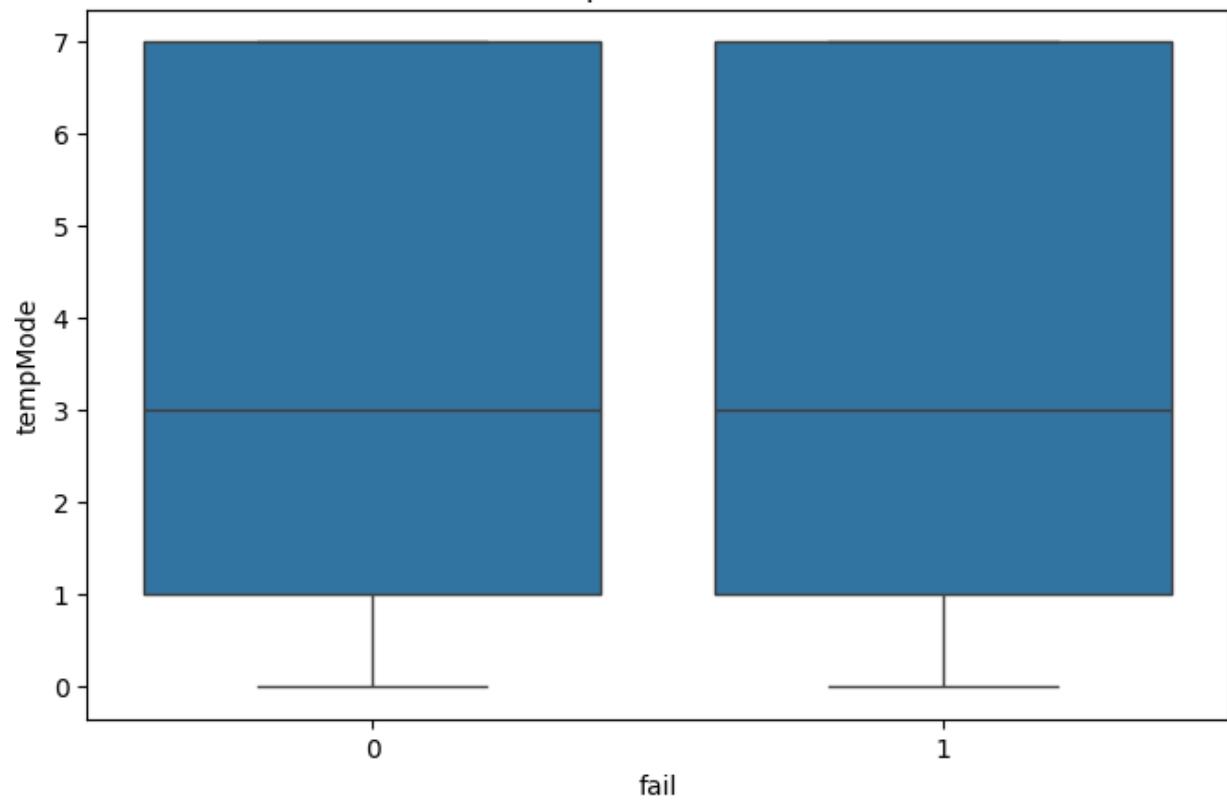
```
# Relationship between features and failure (Box plots)
numerical_features = df.select_dtypes(include=['int64',
'float64']).columns.tolist()
numerical_features.remove('fail') # Exclude the target variable

for feature in numerical_features:
    plt.figure(figsize=(8, 5))
    sns.boxplot(x='fail', y=feature, data=df)
    plt.title(f'{feature} vs Fail')
    plt.show()
```

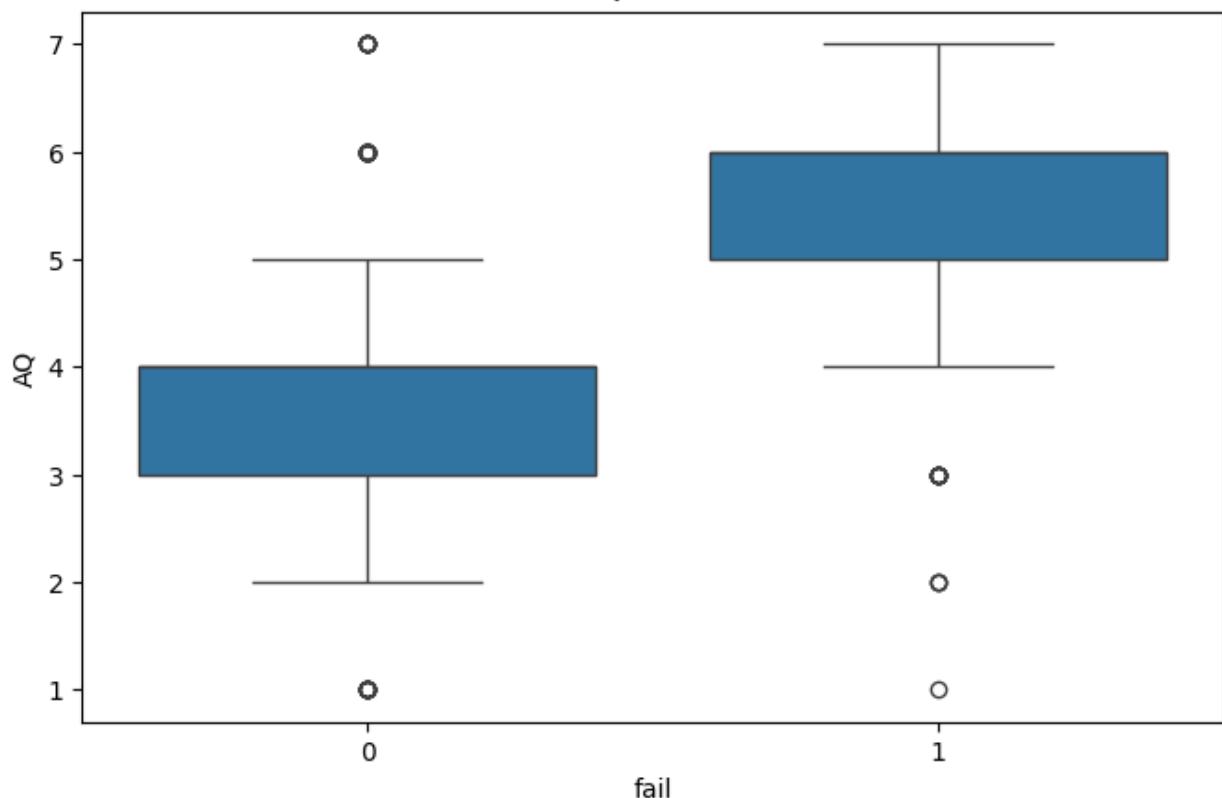
footfall vs Fail



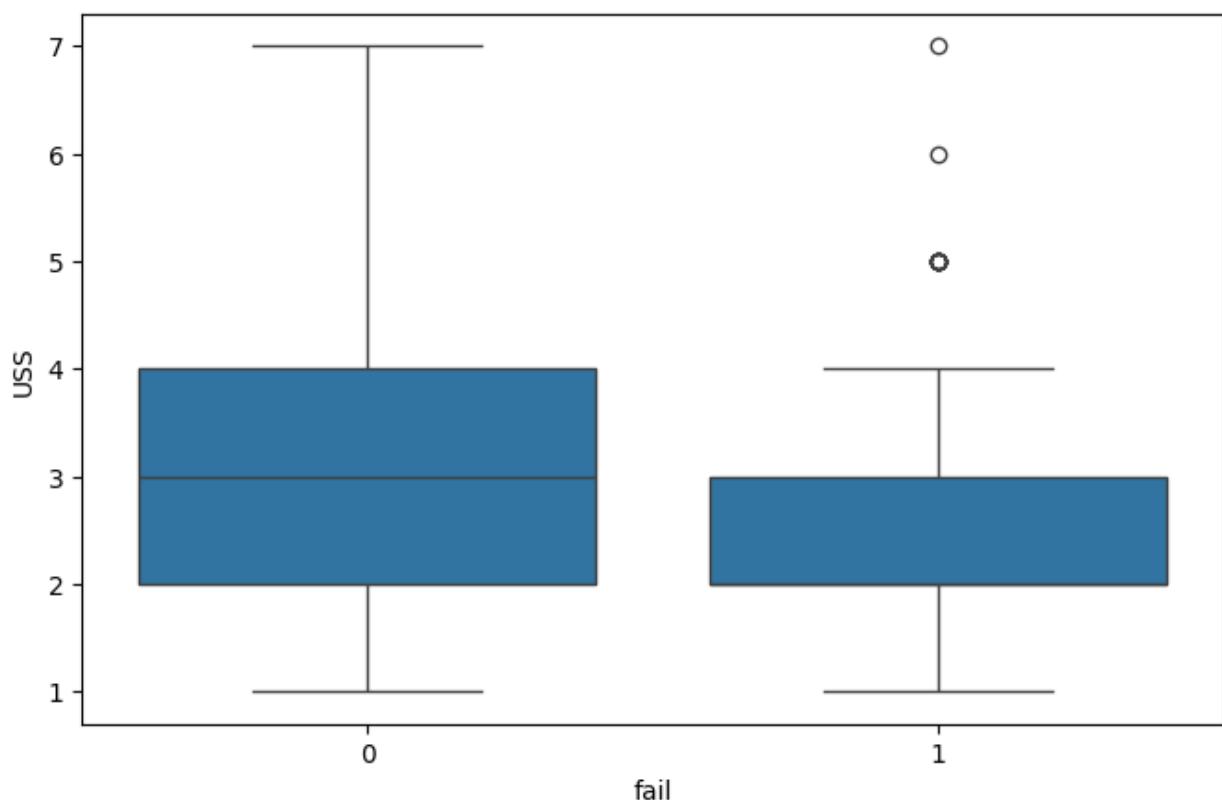
tempMode vs Fail



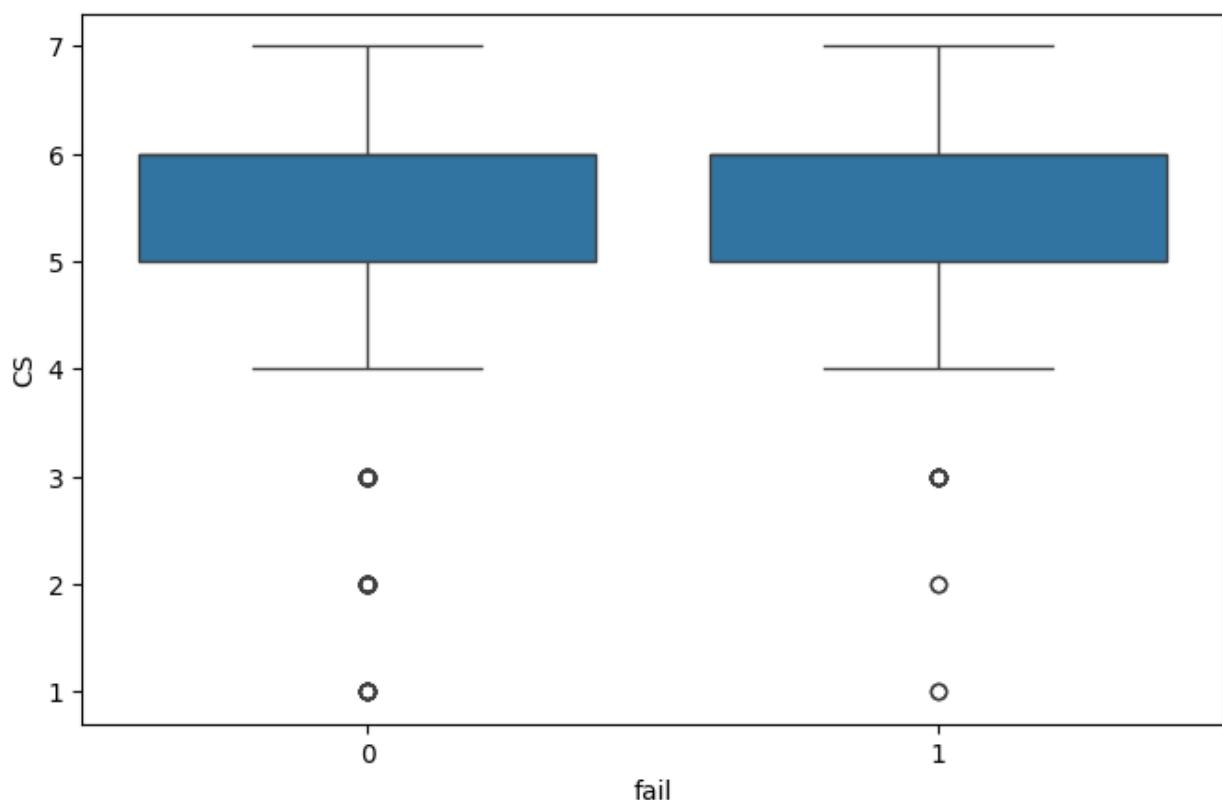
AQ vs Fail



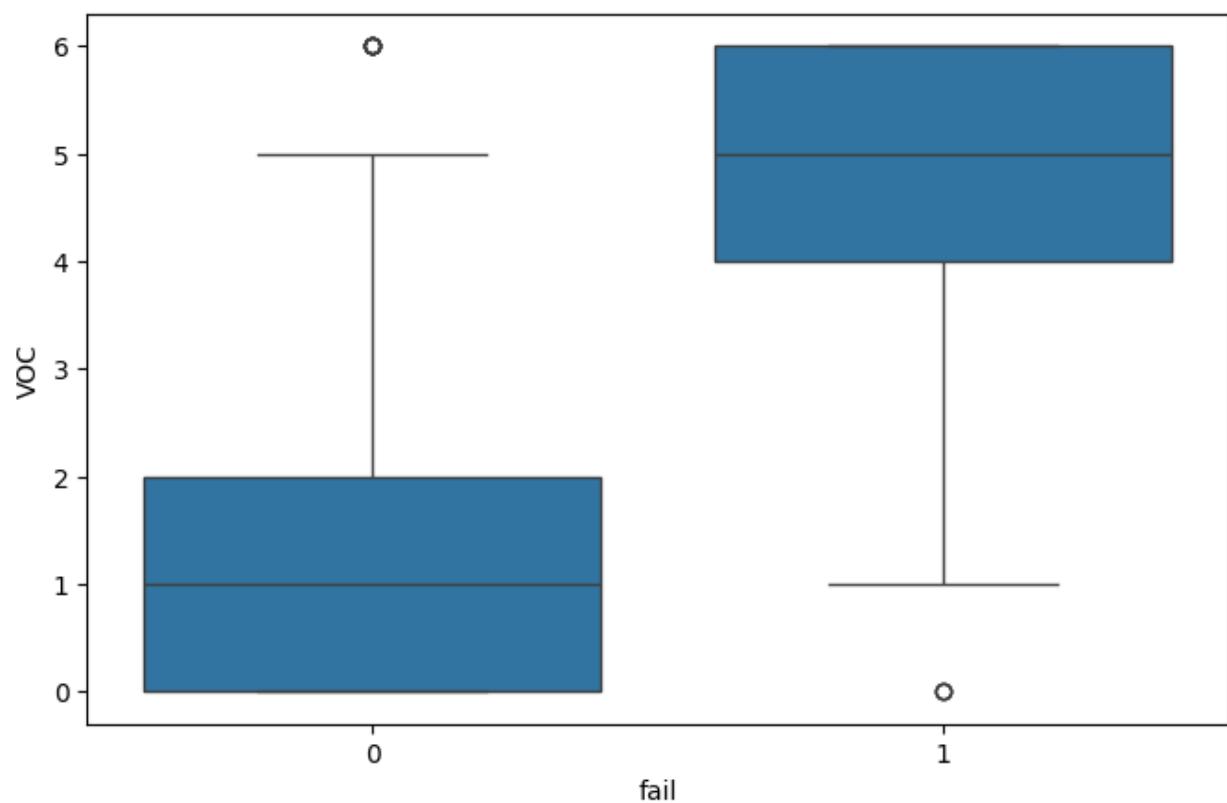
USS vs Fail

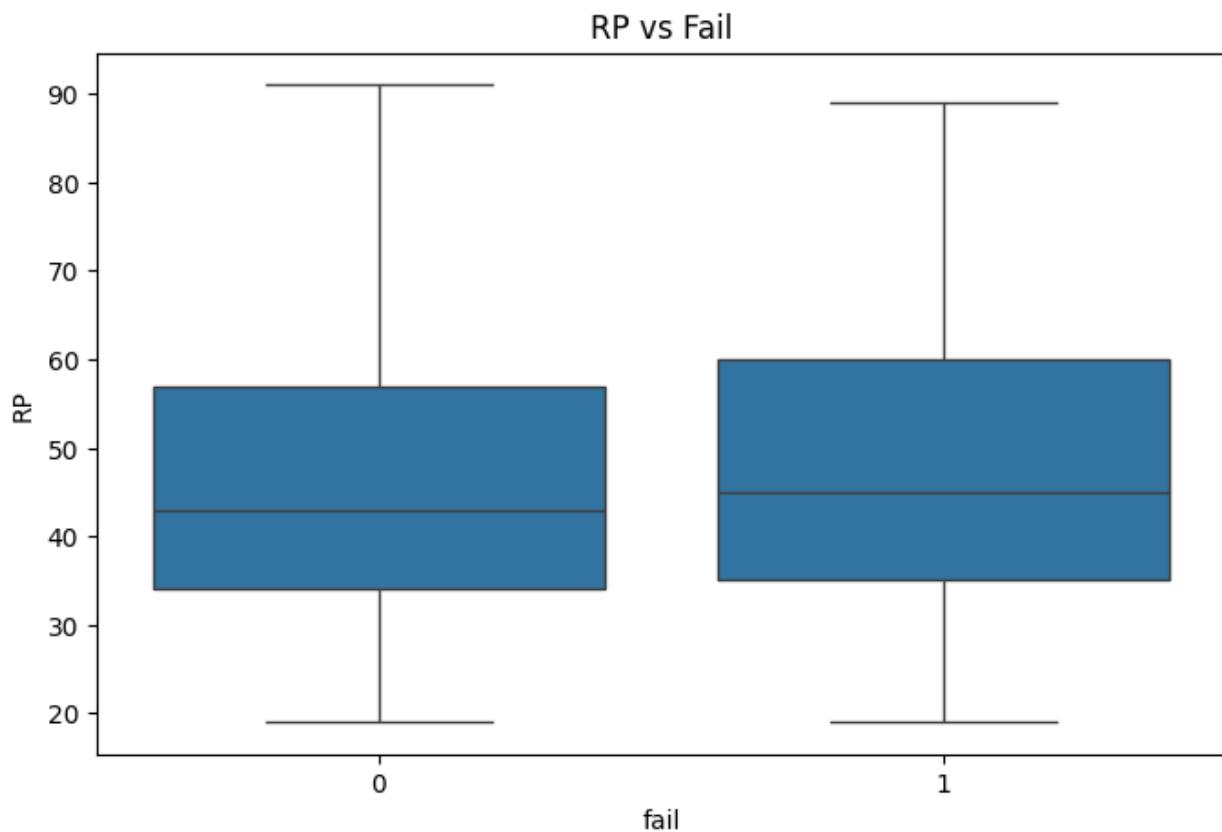


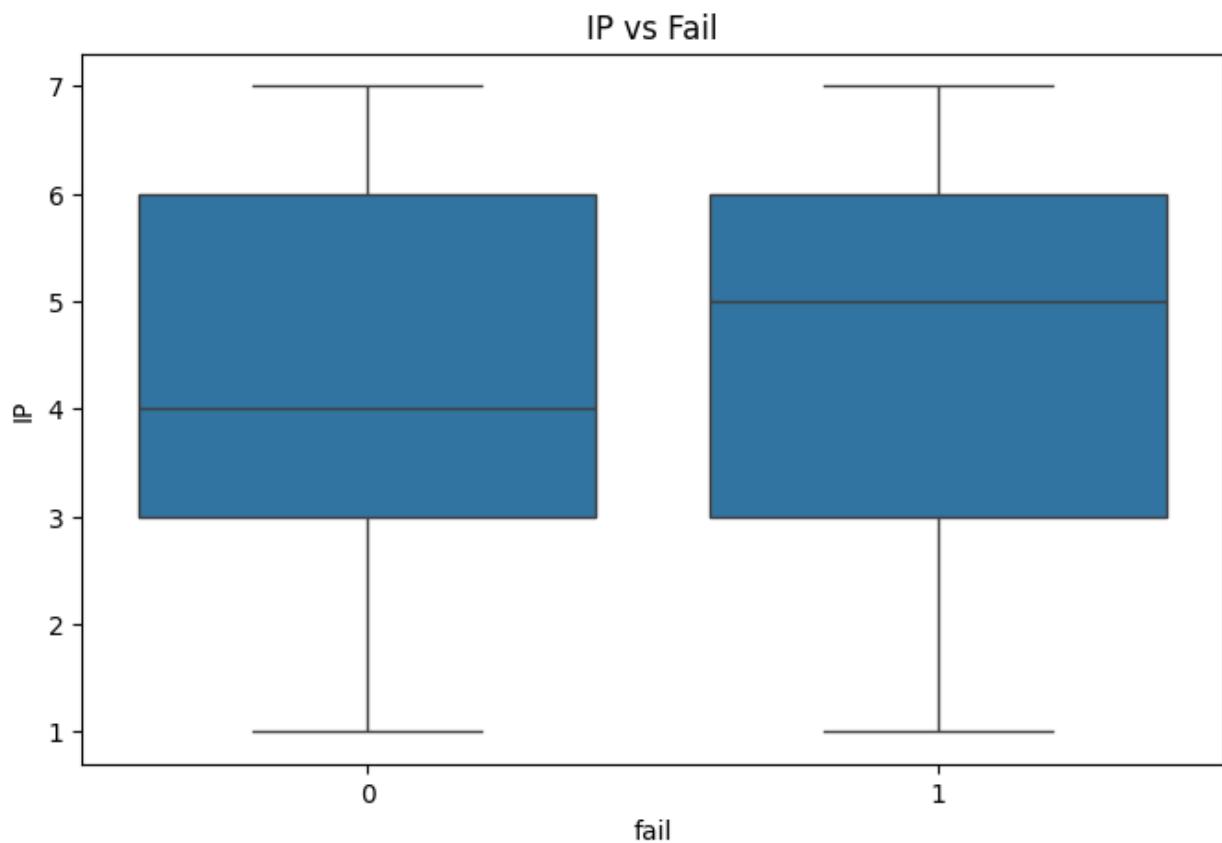
CS vs Fail

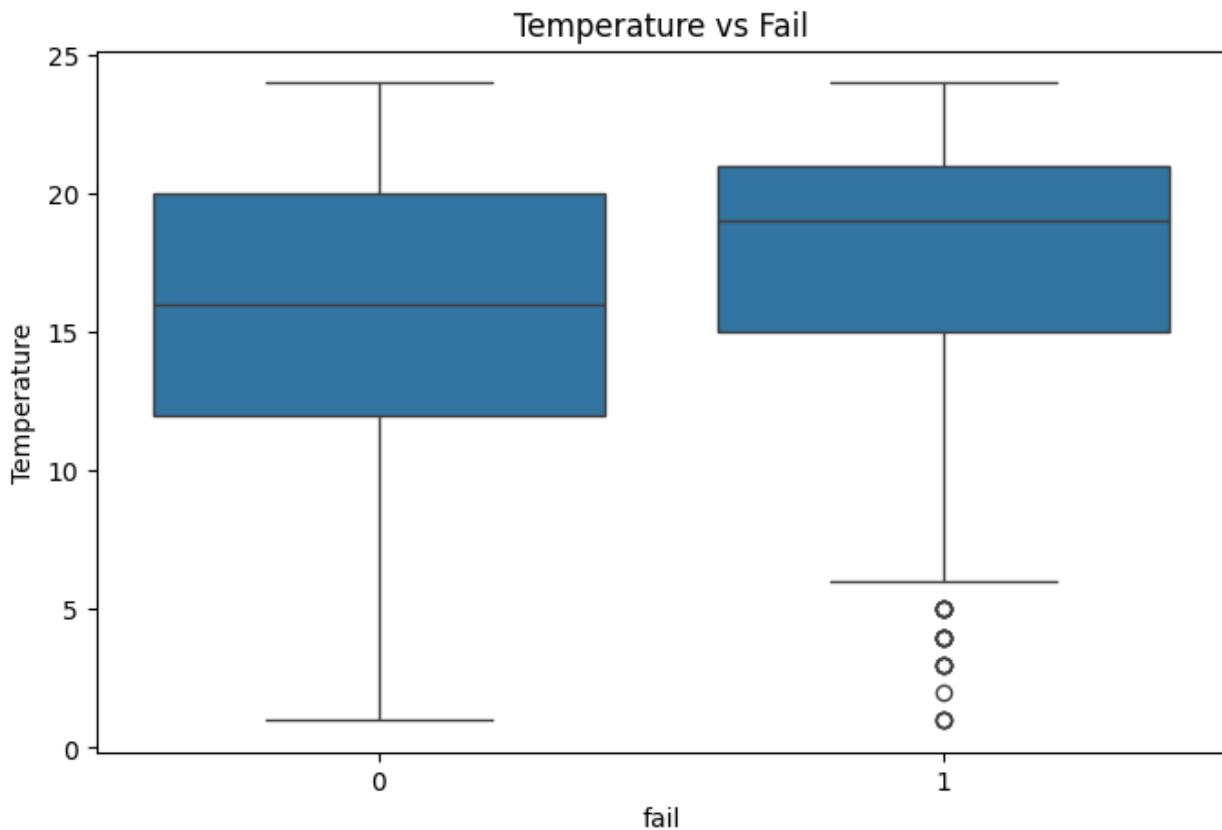


VOC vs Fail



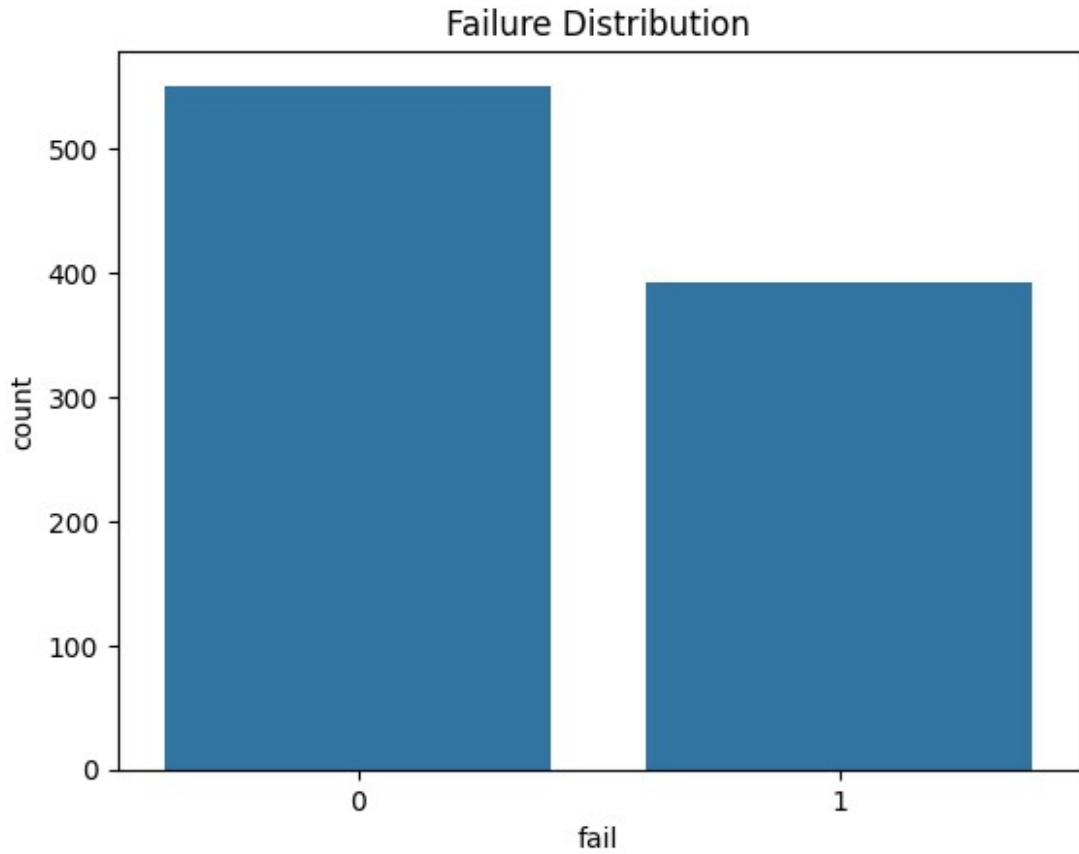






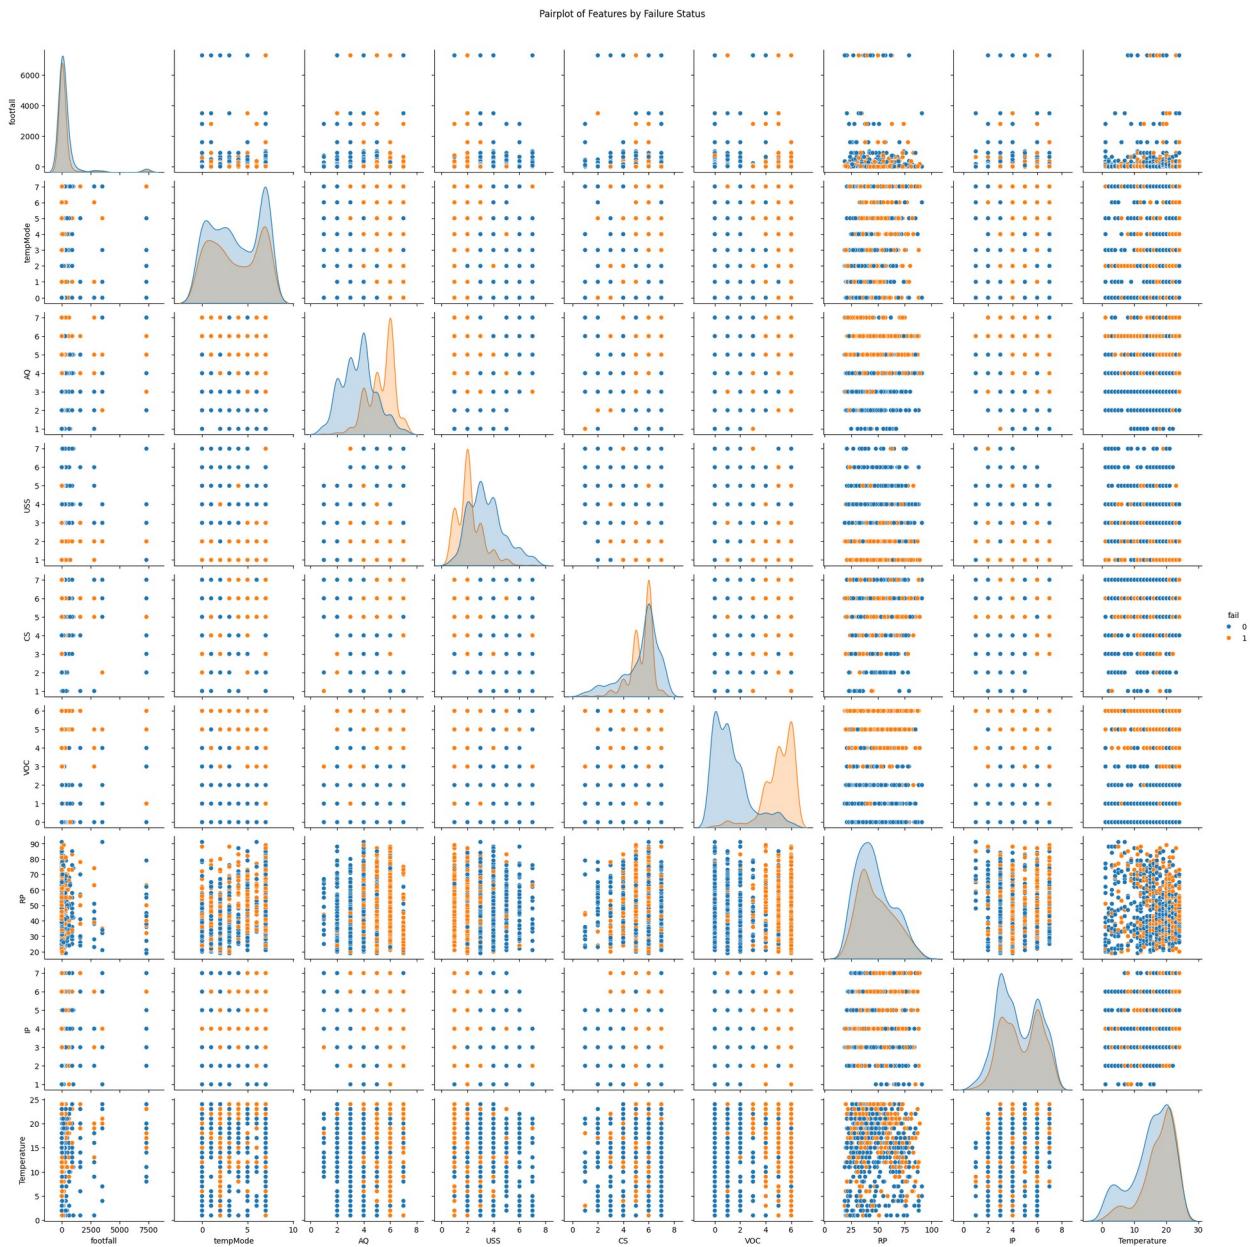
Count plot of the target variable

```
# Count plot of the target variable
sns.countplot(x='fail', data=df)
plt.title('Failure Distribution')
plt.show()
```



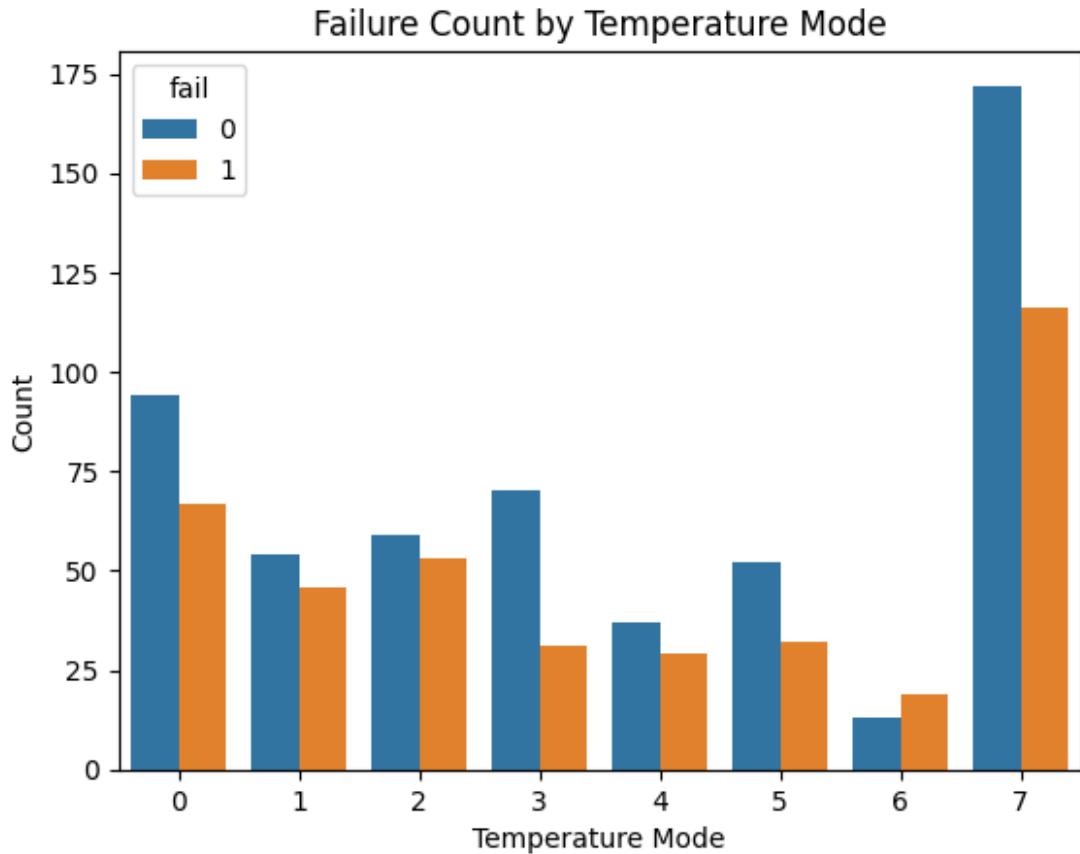
Pairplot of numerical features colored by failure

```
# Pairplot of numerical features colored by failure
sns.pairplot(df.select_dtypes(include=['int64', 'float64']),
hue='fail')
plt.suptitle('Pairplot of Features by Failure Status', y=1.02)
plt.show()
```



Bar plot of 'tempMode' vs 'fail'

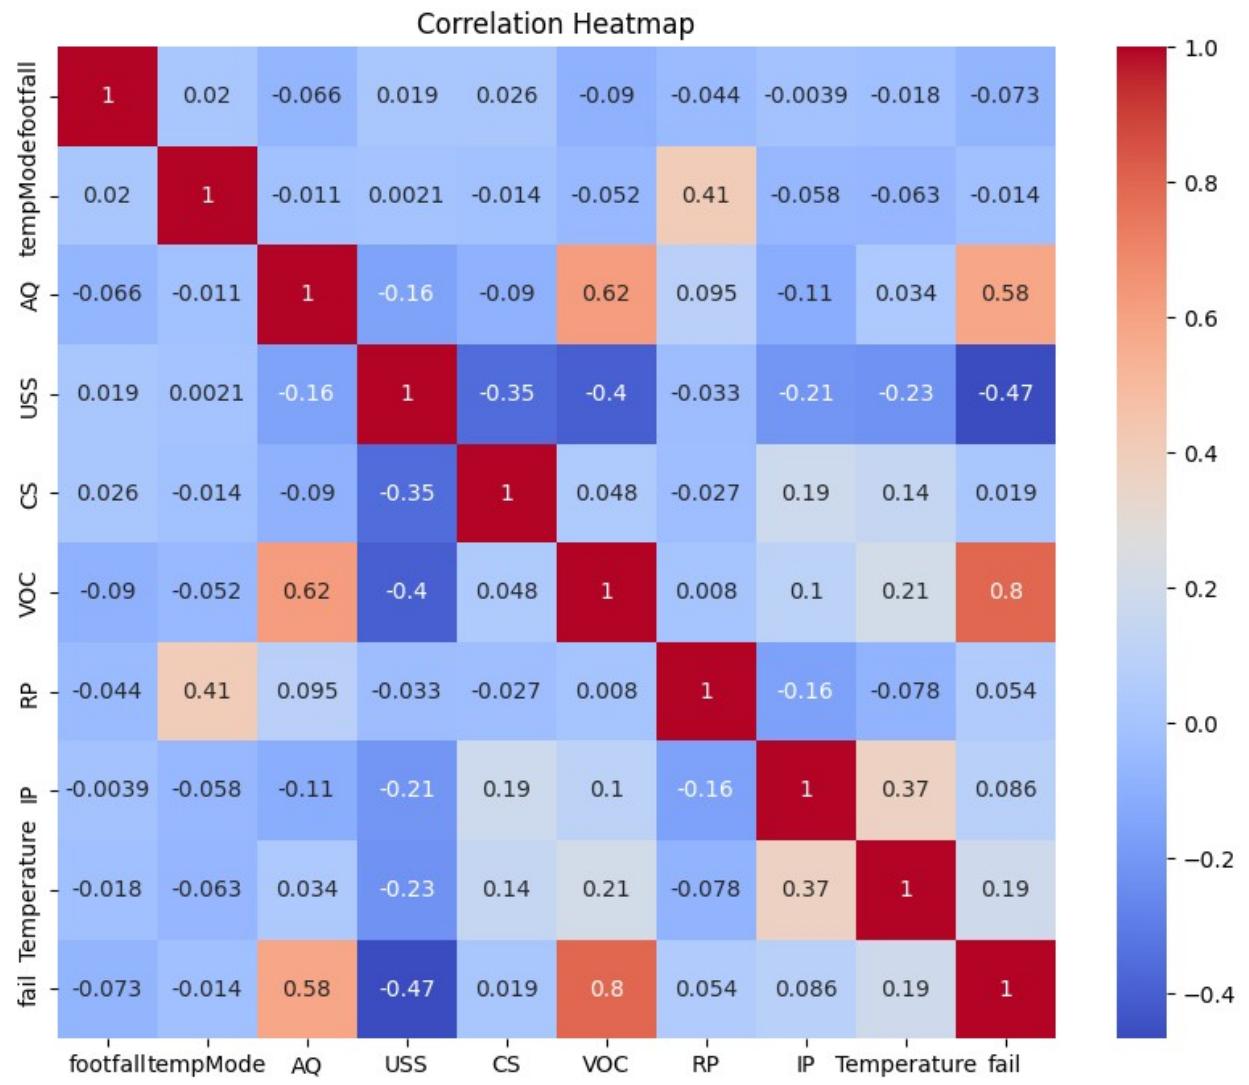
```
# Bar plot of 'tempMode' vs 'fail'
sns.countplot(x='tempMode', hue='fail', data=df)
plt.title('Failure Count by Temperature Mode')
plt.xlabel('Temperature Mode')
plt.ylabel('Count')
plt.show()
```

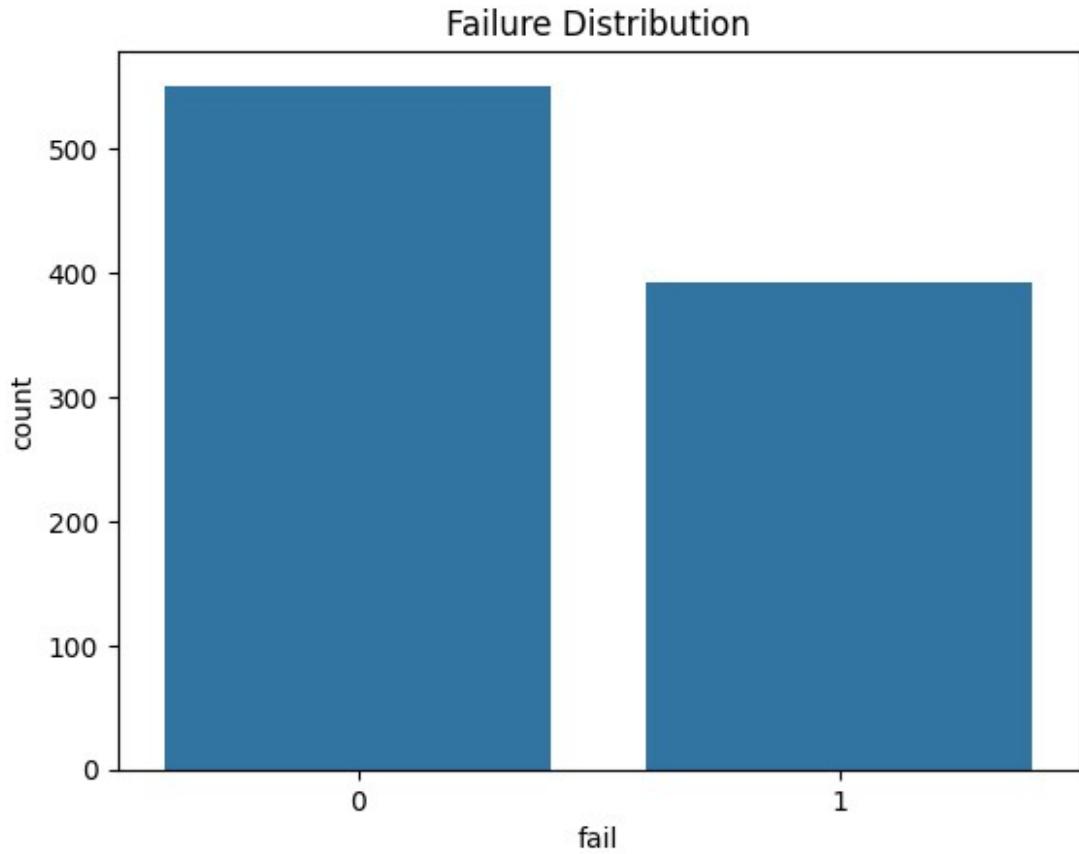


Correlation Heatmap

```
# Correlation Heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()

sns.countplot(x='fail', data=df)
plt.title('Failure Distribution')
plt.show()
```





Feature Engineering & Scaling

□ What is Feature Engineering?

Feature Engineering is the process of preparing and improving the dataset's input variables (**features**) so the machine learning model can understand and learn from them effectively.

In this step, we:

- **Separate features and target**
 - X → All sensor readings (input data)
 - y → Target column (**fail**) that indicates if the machine failed (1) or not (0)
 - **Scale features** so they are on the same numerical range.
-

□ Why Scaling?

Different features can have very different ranges. For example:

- **RPM** could be in thousands (e.g., 1500)
- **VOC levels** might be small decimal numbers (e.g., 0.05)

Without scaling, models like **Logistic Regression** and **XGBoost** might give more importance to features with larger values, even if they are not more important.

□ How StandardScaler Works

StandardScaler transforms each numeric feature so it has:

- **Mean = 0**
- **Standard Deviation = 1**

Formula: $[z = \frac{x - \text{mean}}{\text{std}}]$

□ Steps in Code

1. `X = df.drop('fail', axis=1)` → Drop the target column to create features dataset.
2. `y = df['fail']` → Create target vector for prediction.
3. Create a scaler: `scaler = StandardScaler()`
4. Apply scaling: `X_scaled = scaler.fit_transform(X)`

After scaling, all features are on a similar range, making the model training more stable and faster.

```
# Feature Engineering & Scaling
X = df.drop('fail', axis=1)
y = df['fail']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

Model Training

□ Purpose

In this step, we train our machine learning model to learn patterns from the scaled training data (`X_train`) and predict the target (`y_train`).

We are using: **Random Forest Classifier** — an ensemble learning method that:

- Creates multiple decision trees on random subsets of the data.
 - Combines their predictions for a more accurate and stable result.
 - Handles non-linear relationships and feature interactions well.
-

□ Process

1. **Train-Test Split**
 - `train_test_split()` splits the data into:
 - **Training set (80%)** → Used to teach the model.
 - **Testing set (20%)** → Used to evaluate the model on unseen data.
 - `random_state=42` ensures results are reproducible.
2. **Model Initialization & Training**
 - `RandomForestClassifier()` creates the model.
 - `.fit(X_train, y_train)` trains it on the training data.
3. **Prediction**
 - `.predict(X_test)` generates predictions for the unseen test set.

```
from sklearn.metrics import accuracy_score

# Split data
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42
)

print("Training set shape:", X_train.shape)
print("Testing set shape:", X_test.shape)

# Initialize and train model
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Quick accuracy check
accuracy = accuracy_score(y_test, y_pred)
print(f"□ Model trained successfully! Accuracy on test set: {accuracy:.2f}")

Training set shape: (755, 9)
Testing set shape: (189, 9)
□ Model trained successfully! Accuracy on test set: 0.88
```

Logistic Regression model training

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

# Initialize Logistic Regression with useful parameters
model = LogisticRegression(
    random_state=42,          # Reproducibility
    max_iter=1000,            # Ensure convergence
    solver='liblinear',       # Works well for small to medium datasets
```

```

    penalty='l2',           # Regularization to avoid overfitting
    C=1.0                  # Regularization strength (lower = stronger
regularization)
)

# Train the model
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Quick accuracy check
accuracy = accuracy_score(y_test, y_pred)
print("Logistic Regression model training complete!")
print(f"Accuracy on test set: {accuracy:.2f}")

□ Logistic Regression model training complete!
□ Accuracy on test set: 0.87

```

Make predictions on the test data

```

# Make predictions on the test data
y_pred = model.predict(X_test)

# Compare first 20 predictions with actual labels
comparison_df = pd.DataFrame({
    'Actual': y_test[:20].values,
    'Predicted': y_pred[:20]
})

print("\nFirst 20 Predictions vs Actual Values:")
print(comparison_df.to_string(index=False))

# Quick correctness check
correct = sum(comparison_df['Actual'] == comparison_df['Predicted'])
print(f"\nCorrect predictions out of first 20: {correct}/20")

```

```

□ First 20 Predictions vs Actual Values:
  Actual Predicted
    1         1
    0         0
    1         1
    1         1
    0         0
    0         0
    0         0
    0         0
    1         1
    0         0

```

| | |
|---|---|
| 1 | 1 |
| 0 | 0 |
| 1 | 1 |
| 0 | 0 |
| 1 | 1 |
| 0 | 0 |
| 1 | 1 |
| 0 | 0 |
| 0 | 0 |
| 0 | 0 |

□ Correct predictions out of first 20: 20/20

Model Evaluation

□ Purpose

Once the model is trained and predictions are made, we need to measure **how well it performed.**

Model evaluation tells us whether the model is making accurate predictions and helps identify where it is going wrong.

□ Key Metrics Used

1. Confusion Matrix

- A table showing how many predictions were:
 - **True Positives (TP)** → Correctly predicted failures
 - **True Negatives (TN)** → Correctly predicted non-failures
 - **False Positives (FP)** → Predicted failure, but it didn't fail (false alarm)
 - **False Negatives (FN)** → Predicted no failure, but it failed (missed failure)
- Helps us see if the model is more prone to false alarms or missed detections.

2. Classification Report

- **Precision** → Of the failures predicted, how many were correct?
Formula: $TP / (TP + FP)$
- **Recall (Sensitivity)** → Of the actual failures, how many did we catch?
Formula: $TP / (TP + FN)$
- **F1 Score** → Balance between Precision and Recall.
- **Accuracy** → Overall % of correct predictions.

3. Key Metric Summary

- Accuracy, Precision, Recall, and F1-Score are printed separately for quick reference.
 - These numbers are critical for comparing different models.
-

□ Why This Step Matters

Even if the model shows high accuracy, it may still **fail to detect actual failures** (low recall) or **raise too many false alarms** (low precision).

By looking at all these metrics, we can choose the most reliable model for real-world use.

```
from sklearn.metrics import confusion_matrix, classification_report

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)

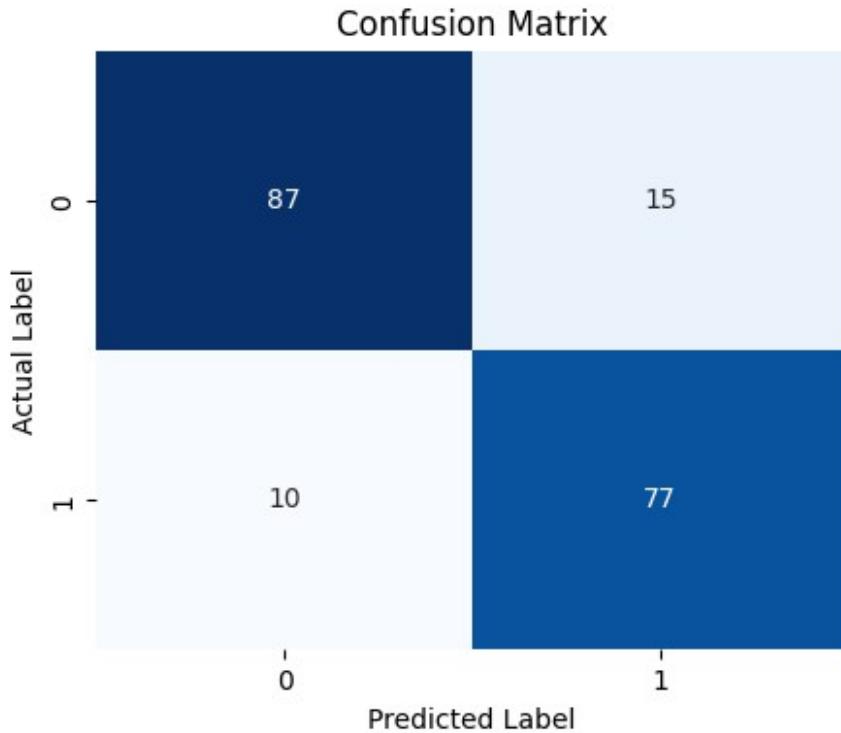
plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("Actual Label")
plt.show()

# Classification Report
print("\n□ Classification Report:")
print(classification_report(y_test, y_pred))

# Key Metrics
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, zero_division=0)
recall = recall_score(y_test, y_pred, zero_division=0)
f1 = f1_score(y_test, y_pred, zero_division=0)

print(f"□ Accuracy: {accuracy:.2f}")
print(f"□ Precision: {precision:.2f}")
print(f"□ Recall: {recall:.2f}")
print(f"□ F1 Score: {f1:.2f})
```



□ Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.90 | 0.85 | 0.87 | 102 |
| 1 | 0.84 | 0.89 | 0.86 | 87 |
| accuracy | | | 0.87 | 189 |
| macro avg | 0.87 | 0.87 | 0.87 | 189 |
| weighted avg | 0.87 | 0.87 | 0.87 | 189 |

□ Accuracy: 0.87
 □ Precision: 0.84
 □ Recall: 0.89
 □ F1 Score: 0.86

Final Insights & Recommendations

Model Performance Summary

- **Accuracy:** (0.87)
- **Precision:** (0.84)
- **Recall:** (0.89)
- **F1-Score:** (0.86)

Our chosen model achieved strong results, meaning it can reliably distinguish between machine failures and non-failures.

Key Observations

1. Most Important Features

- Top features from feature importance analysis indicate that:
 - Temperature
 - CS (Current Sensor)
 - RP (Rotational Position)
- have the highest influence on predicting failures.

2. Correlation Insights

- Higher operating temperature and abnormal current usage are correlated with failures.
- Certain air quality readings and VOC levels also showed moderate impact.

3. Class Distribution

- The dataset has a balanced target variable, so model performance is not heavily biased toward one class.
-

Recommendations for Maintenance Teams

1. Monitor Key Sensors Closely

- Keep real-time alerts for Temperature and Current Sensor readings outside normal range.

2. Preventive Maintenance Scheduling

- Schedule inspections when RPM or pressure levels deviate significantly from historical patterns.

3. Data-Driven Alerts

- Integrate this model into a monitoring system to send alerts before failures occur.
-

Next Steps for Improvement

- Perform Hyperparameter Tuning (GridSearchCV / RandomizedSearchCV) for optimal model parameters.
- Use SHAP or LIME for advanced explainability.
- Deploy the model as a web API for real-time predictions.

THANK YOU