

🔧 Machine Failure Prediction - Data Analytics Capstone Project

This notebook walks through a full machine learning pipeline to predict machine failures using sensor data.

▼ Data Overview

Project Name: *Data Analytics Major Project – Predictive Maintenance using Sensor Data*

Objective: To predict machine failures in advance using historical sensor readings from various machines, enabling proactive maintenance and minimizing downtime.

Dataset Summary: The dataset consists of multiple sensor readings, environmental parameters, and operational metrics from machines, along with a binary indicator of whether a failure occurred.

Number of Features: 11 (including the target variable) **Target Variable:** Fail (1 = Failure, 0 = No failure)

Feature Descriptions

Feature	Description
Footfall	Number of people or objects passing near the machine.
Temp Mode	Temperature mode/setting of the machine.
AQ	Air Quality Index near the machine.
USS	Ultrasonic sensor readings (proximity measurements).
CS	Electrical current usage (current sensor readings).
VOC	Volatile organic compounds level near the machine.
RP	Rotational position or RPM (revolutions per minute) of machine parts.
IP	Input pressure to the machine.
Temperature	Actual operating temperature of the machine.
Fail	Binary indicator of machine failure (1 = Failure, 0 = No failure).

Potential Use Cases

- **Predictive Maintenance:** Forecast failures before they occur to schedule timely servicing.
- **Operational Efficiency:** Optimize operating conditions to extend machine life.
- **Safety Monitoring:** Detect unsafe conditions based on environmental and sensor readings.

- **Cost Reduction:** Minimize unexpected breakdowns and repair costs.

Let's Begin

▼ Import Libraries

```
# Step 1: Import Libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix
import warnings
warnings.filterwarnings('ignore')
```

▼ Load Dataset

```
# Step 2: Load Dataset
df = pd.read_csv('/content/data (1).csv')
```

▼ Dataset First View

```
# Dataset First Look
df.head()
```

	footfall	tempMode	AQ	USS	CS	VOC	RP	IP	Temperature	fail	
0	0	7	7	1	6	6	36	3		1	1
1	190	1	3	3	5	1	20	4		1	0
2	31	7	2	2	6	1	24	6		1	0
3	83	4	3	4	5	1	28	6		1	0
4	640	7	5	6	4	0	68	6		1	0

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

Double-click (or enter) to edit

▼ Dataset Rows & Columns count

```
# Dataset Rows & Columns
```

```
df.shape
```

```
→ (944, 10)
```

▼ Dataset Information

```
# Dataset Info
```

```
df.info()
```

```
→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 944 entries, 0 to 943
Data columns (total 10 columns):
 #   Column      Non-Null Count  Dtype  
---  --          --          --      
 0   footfall    944 non-null    int64  
 1   tempMode    944 non-null    int64  
 2   AQ          944 non-null    int64  
 3   USS         944 non-null    int64  
 4   CS          944 non-null    int64  
 5   VOC         944 non-null    int64  
 6   RP          944 non-null    int64  
 7   IP          944 non-null    int64  
 8   Temperature 944 non-null    int64  
 9   fail        944 non-null    int64  
dtypes: int64(10)
memory usage: 73.9 KB
```

▼ Duplicate Values

```
# Dataset Duplicate Value Count
```

```
df.duplicated().sum()
```

```
→ np.int64(1)
```

▼ Missing Values/Null Values

```
# Missing Values/Null Values Count
```

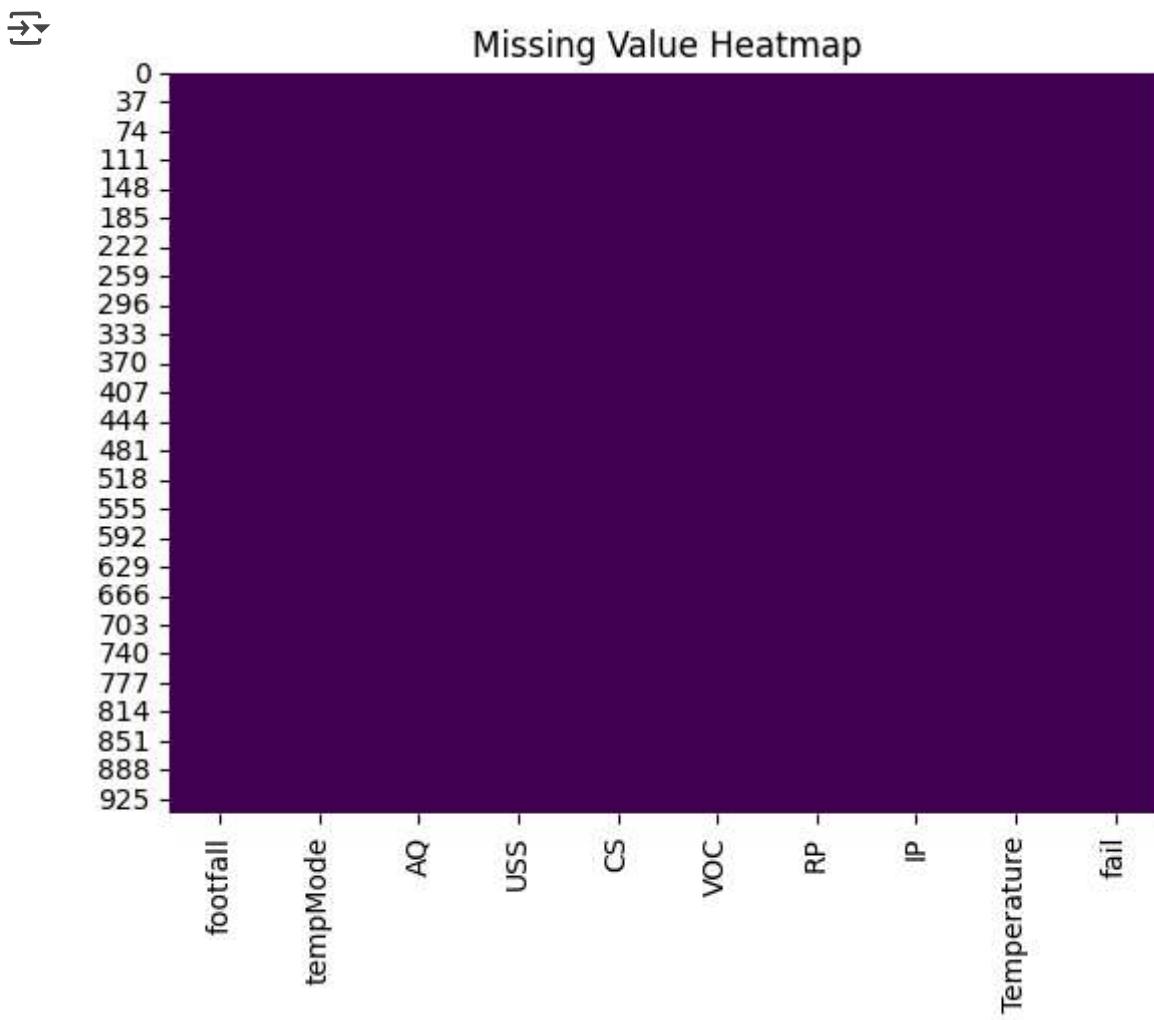
```
df.isnull().sum()
```

	0
footfall	0
tempMode	0
AQ	0
USS	0
CS	0
VOC	0
RP	0
IP	0
Temperature	0
fail	0

dtype: int64

```
# Visualizing the missing values
```

```
sns.heatmap(df.isnull(), cbar=False, cmap='viridis', xticklabels=df.columns)
plt.title('Missing Value Heatmap')
plt.show()
```



▼ Data Cleaning

```
# Step 3: Data Cleaning
print(df.info())
print(df.isnull().sum())
# Fill missing numeric values with median
df.fillna(df.median(), inplace=True)
```

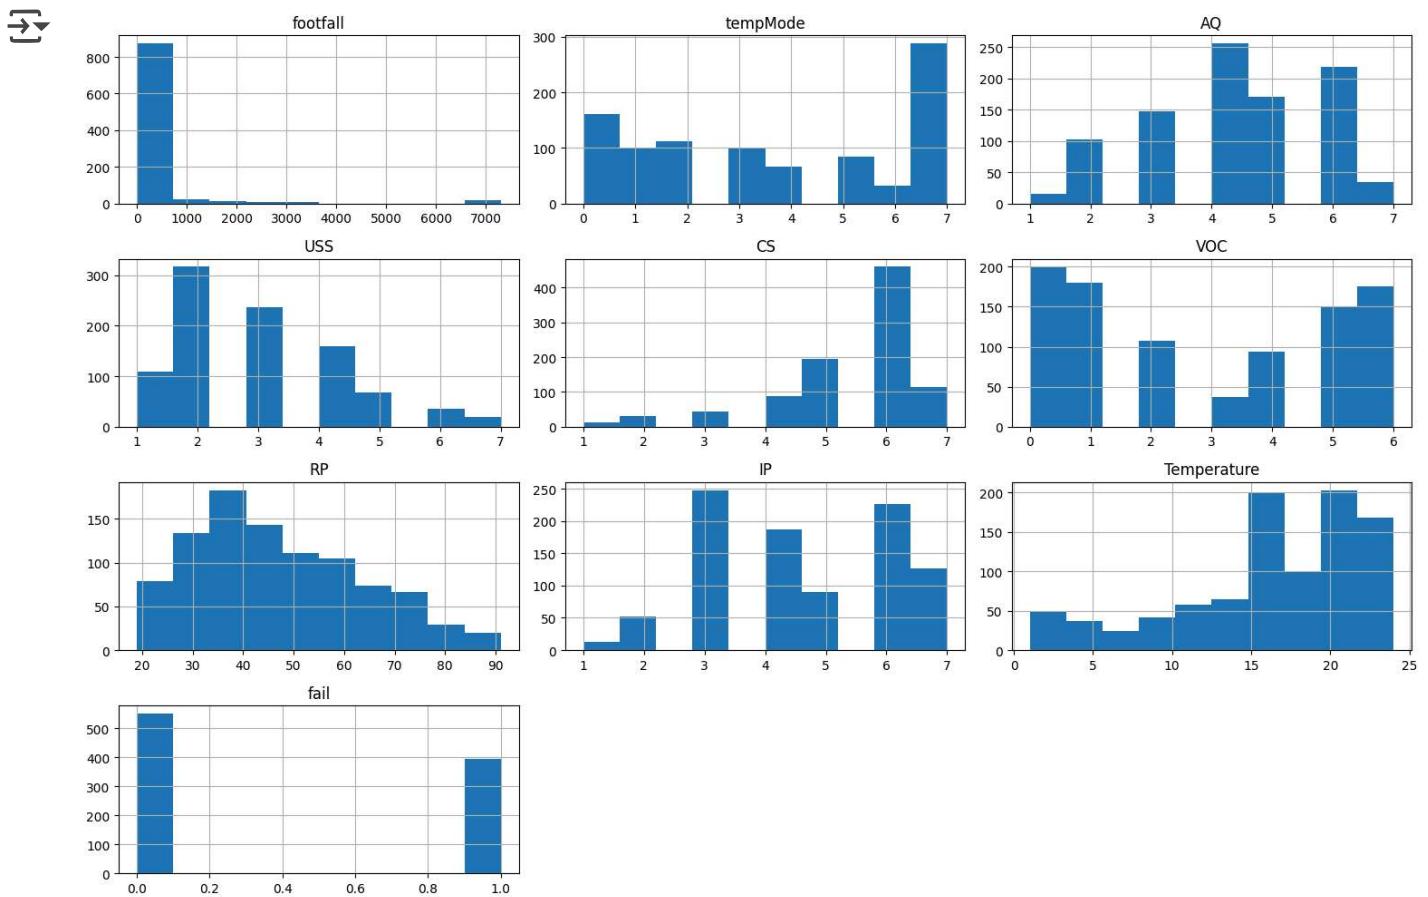
```
→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 944 entries, 0 to 943
Data columns (total 10 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   footfall    944 non-null    int64  
 1   tempMode    944 non-null    int64  
 2   AQ          944 non-null    int64  
 3   USS         944 non-null    int64  
 4   CS          944 non-null    int64  
 5   VOC         944 non-null    int64  
 6   RP          944 non-null    int64  
 7   IP          944 non-null    int64  
 8   Temperature 944 non-null    float64
 9   fail        944 non-null    float64
```

```
8    Temperature    944 non-null    int64
9    fail          944 non-null    int64
dtypes: int64(10)
memory usage: 73.9 KB
None
footfall      0
tempMode       0
AQ            0
USS           0
CS            0
VOC           0
RP            0
IP            0
Temperature   0
fail          0
dtype: int64
```

▼ Data Visualisation

▼ Distribution of each numerical feature

```
# Visualize the distribution of each numerical feature
df.hist(figsize=(15, 10))
plt.tight_layout()
plt.show()
```



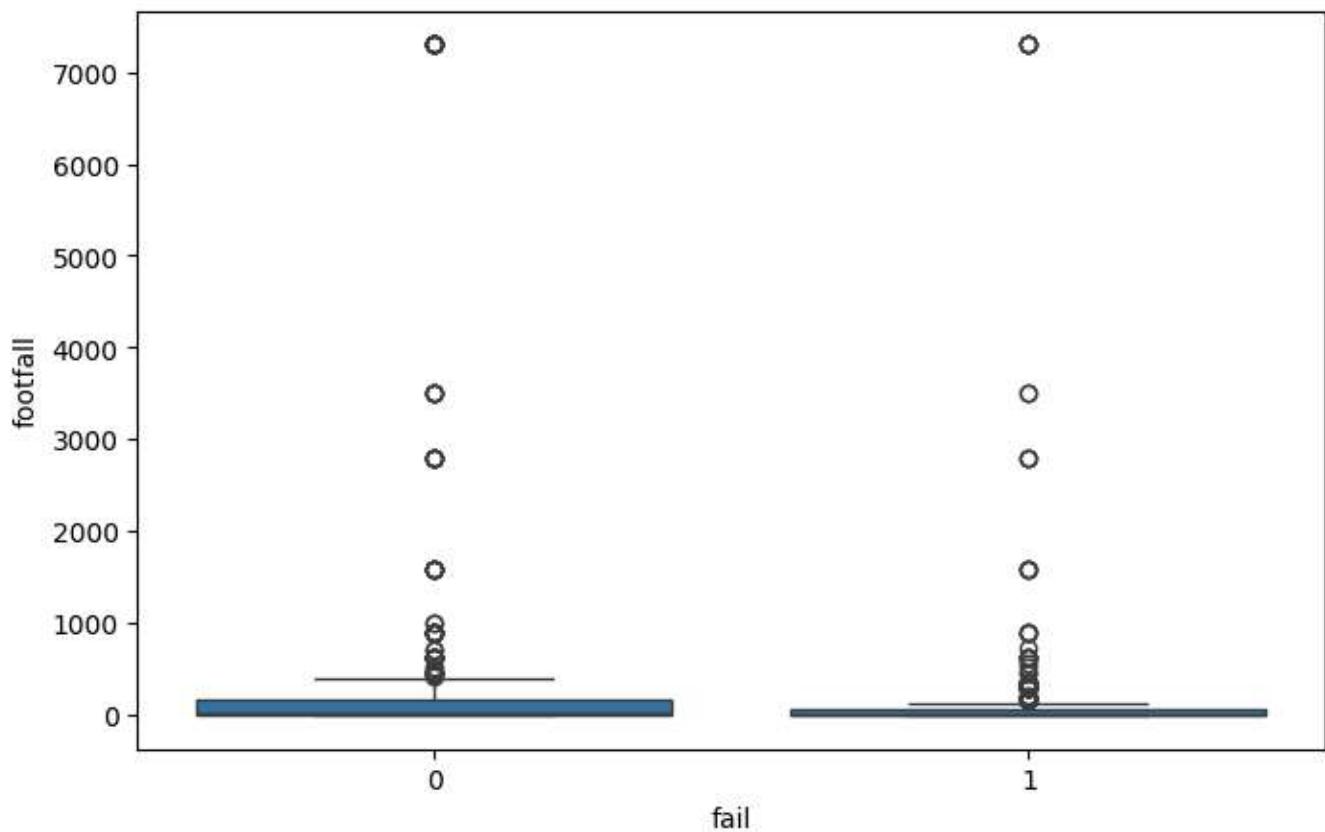
❖ Relationship between features and failure (Box plots)

```
# Relationship between features and failure (Box plots)
numerical_features = df.select_dtypes(include=['int64', 'float64']).columns.tolist()
numerical_features.remove('fail') # Exclude the target variable

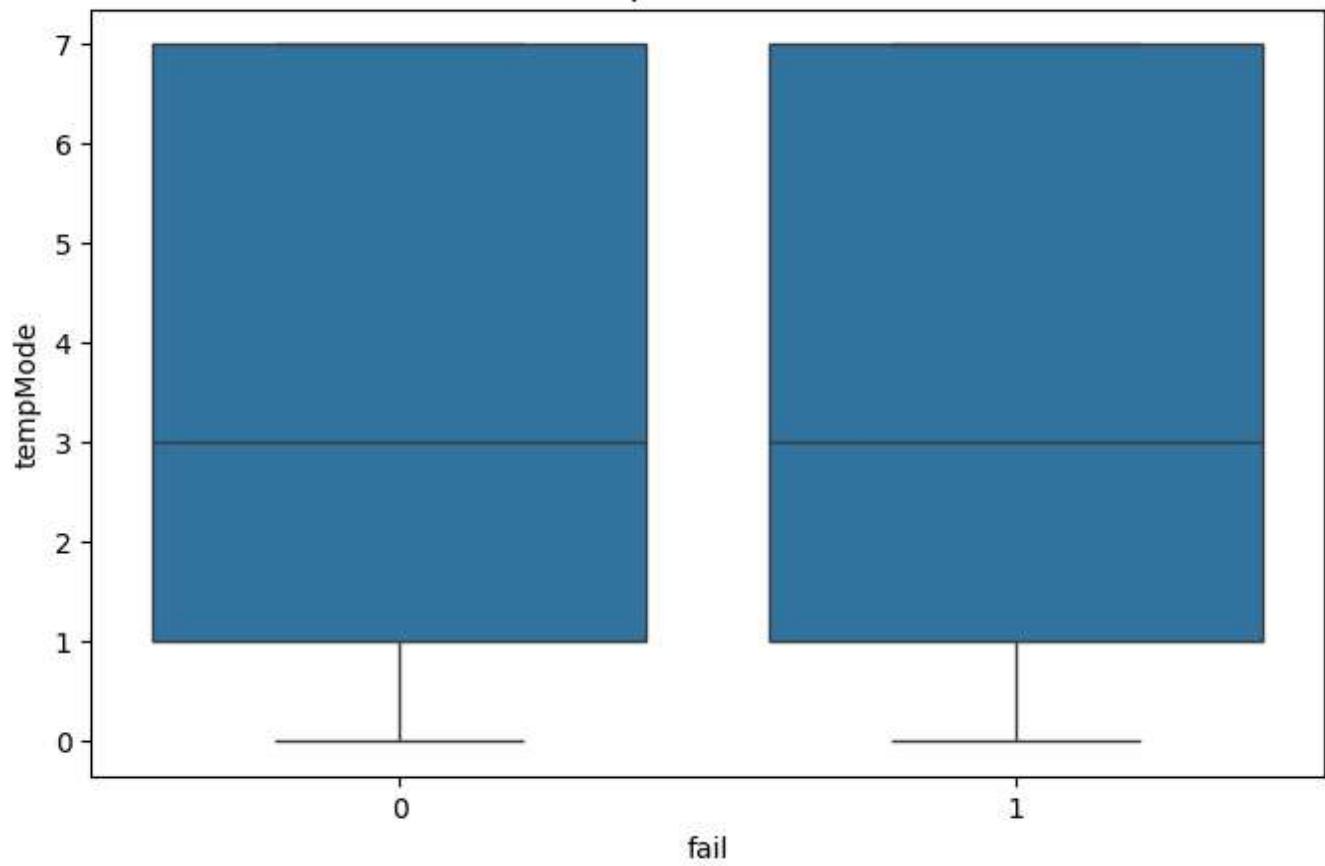
for feature in numerical_features:
    plt.figure(figsize=(8, 5))
    sns.boxplot(x='fail', y=feature, data=df)
    plt.title(f'{feature} vs Fail')
    plt.show()
```



footfall vs Fail

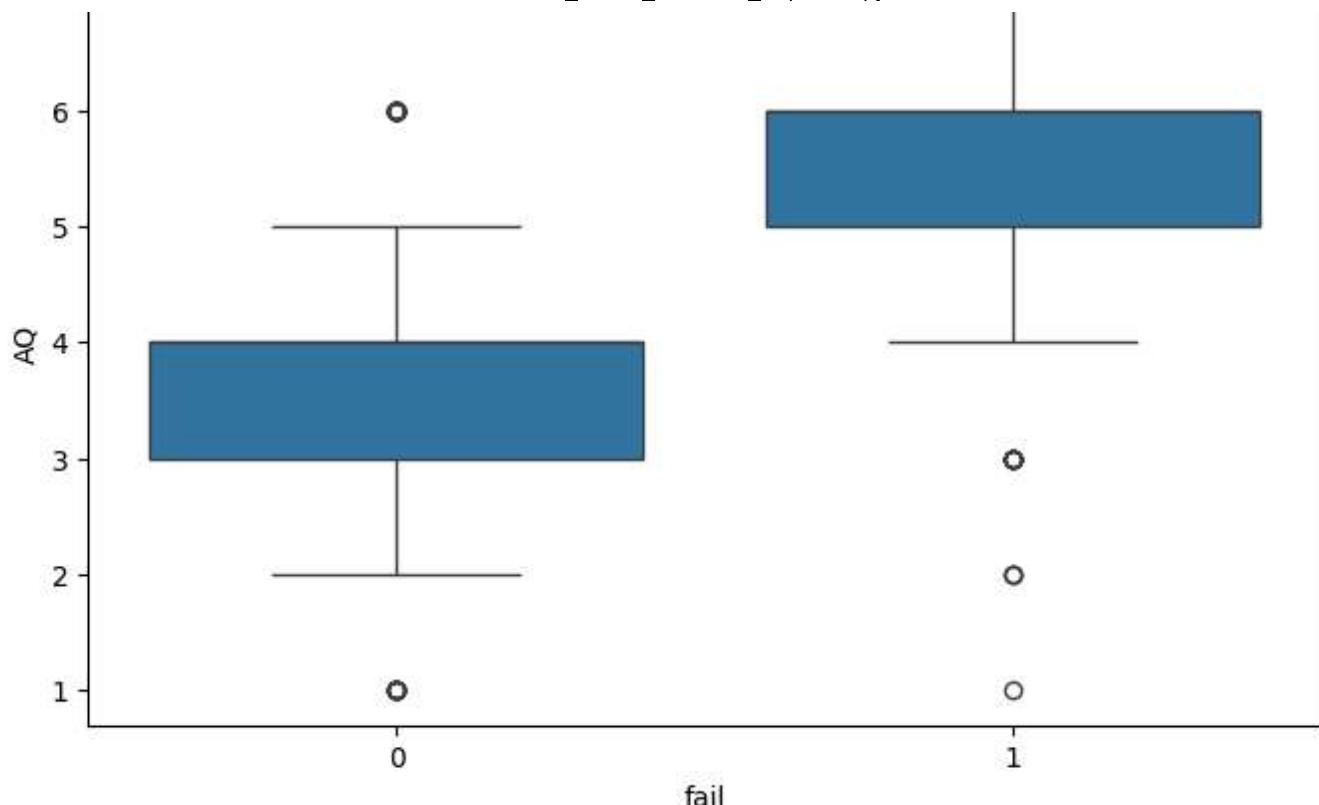


tempMode vs Fail

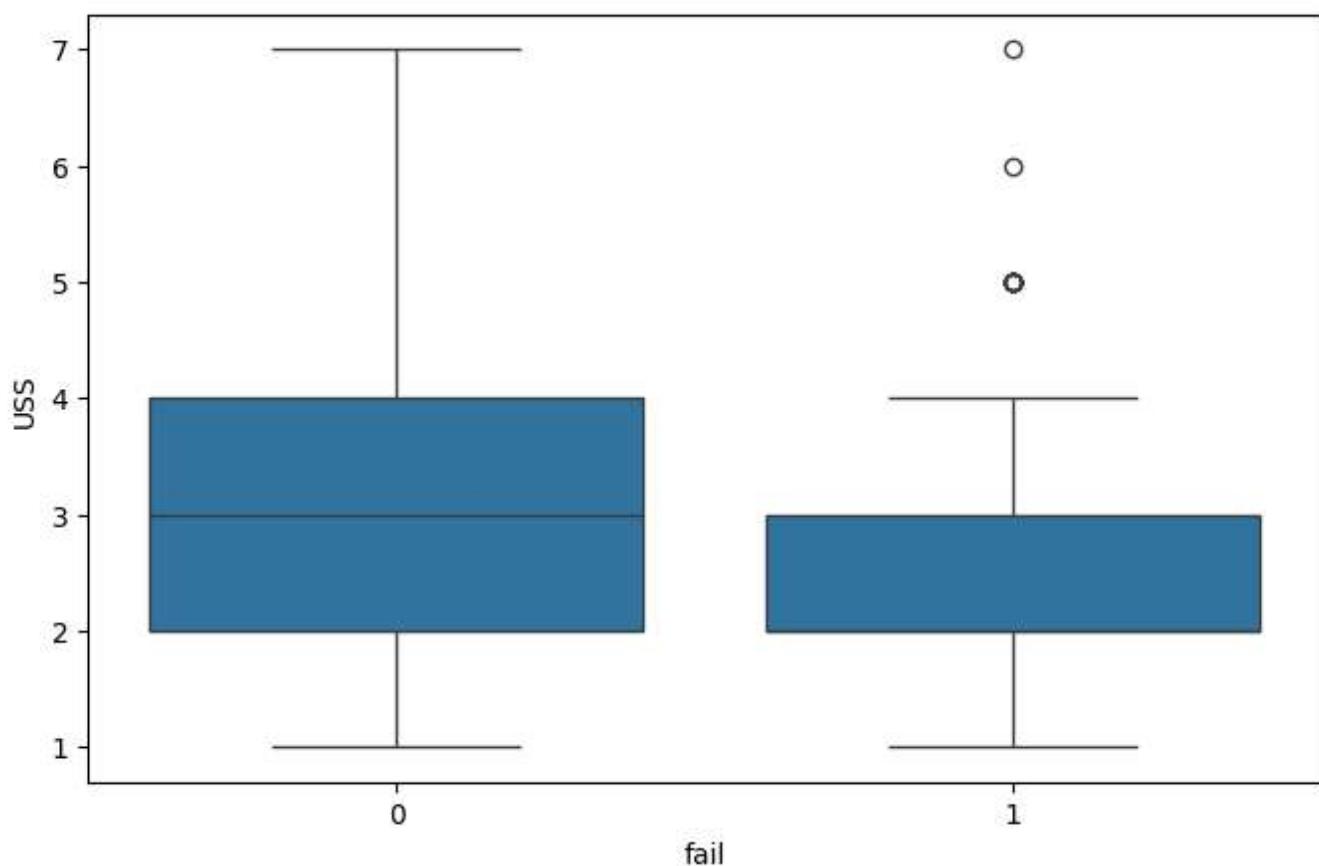


AQ vs Fail

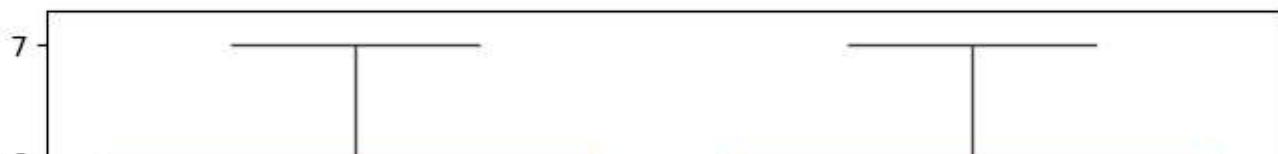


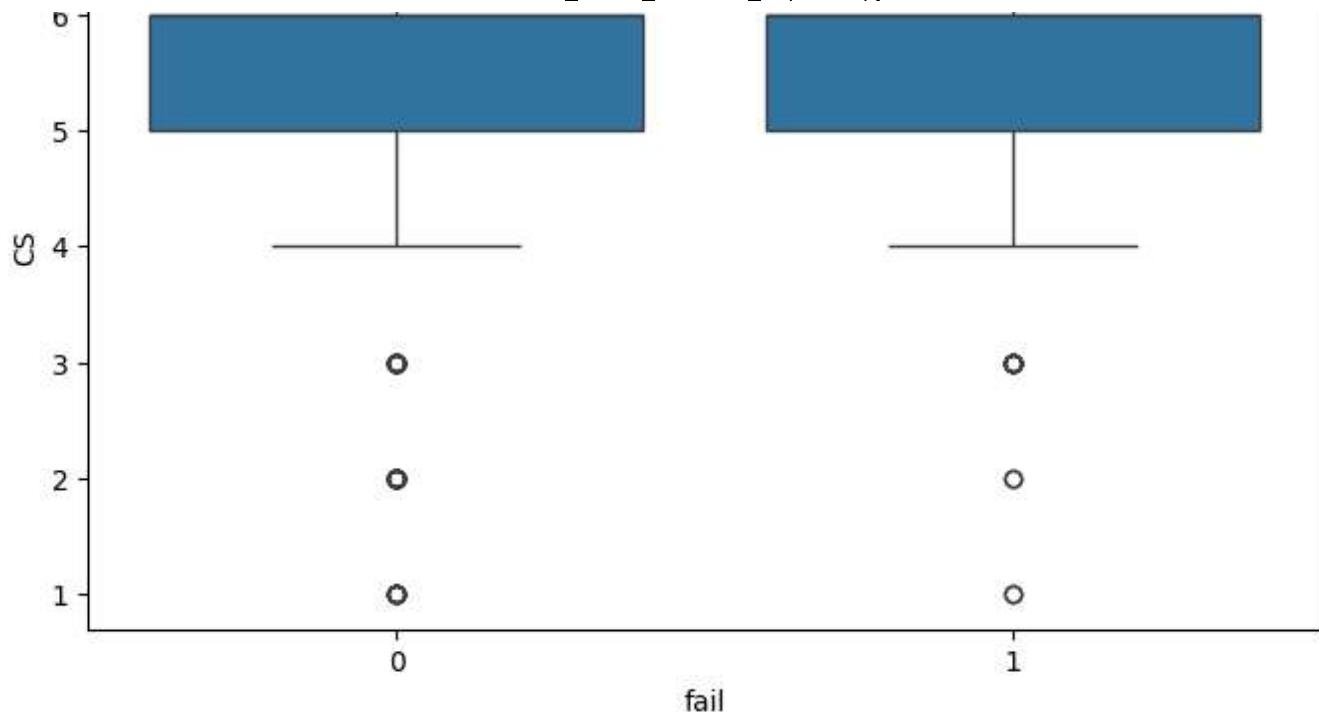


USS vs Fail

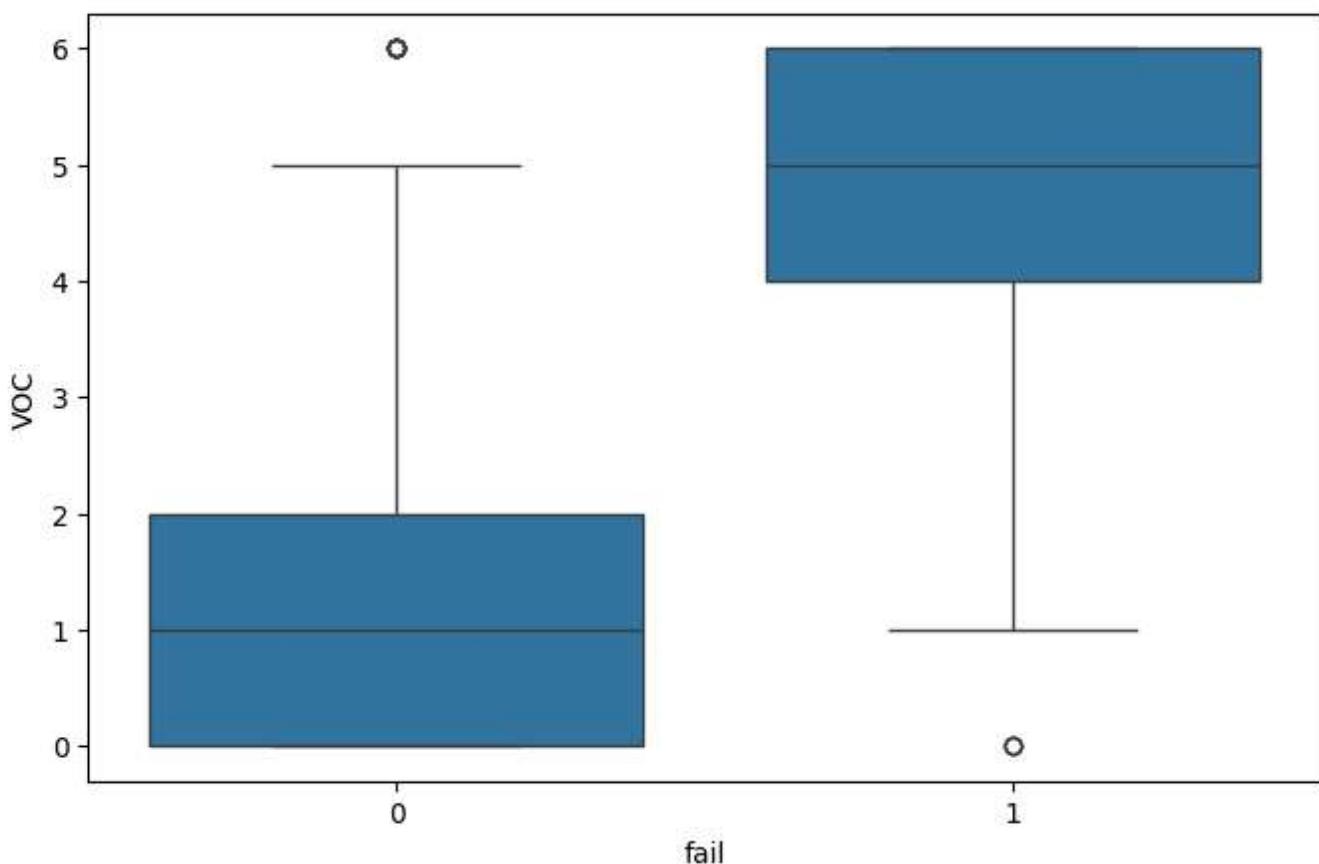


CS vs Fail

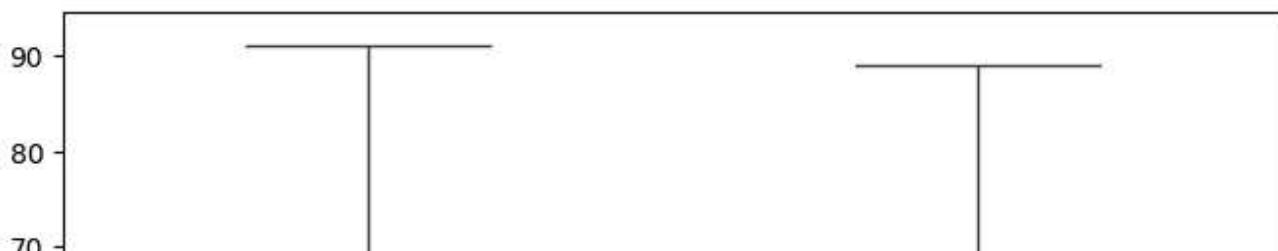


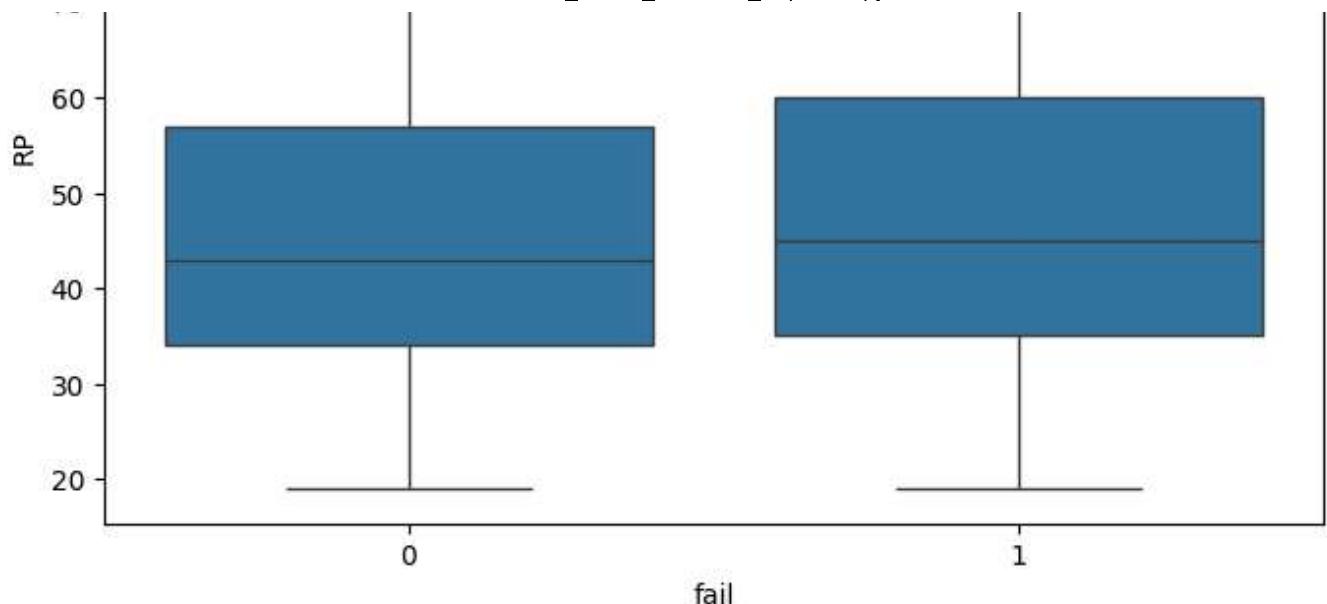


VOC vs Fail

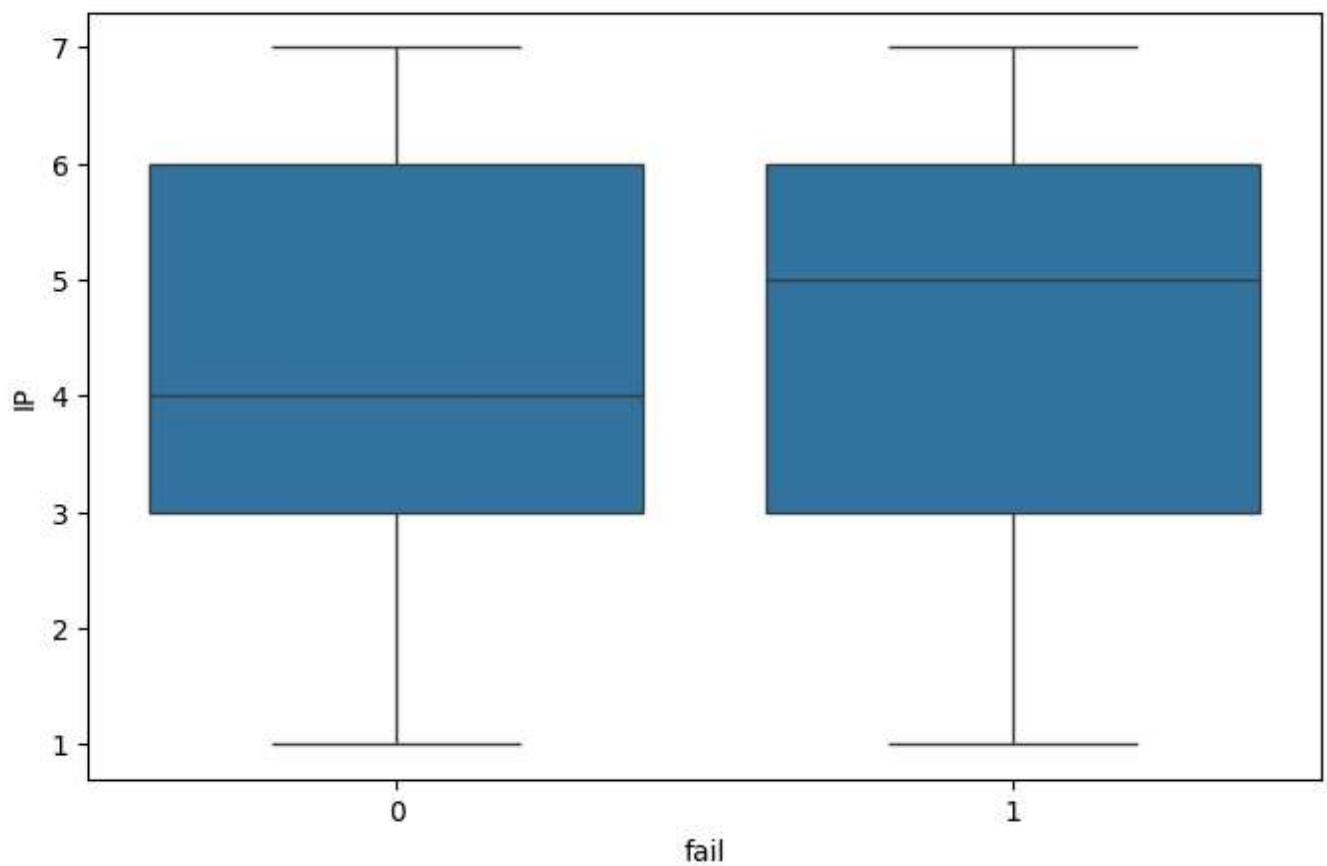


RP vs Fail

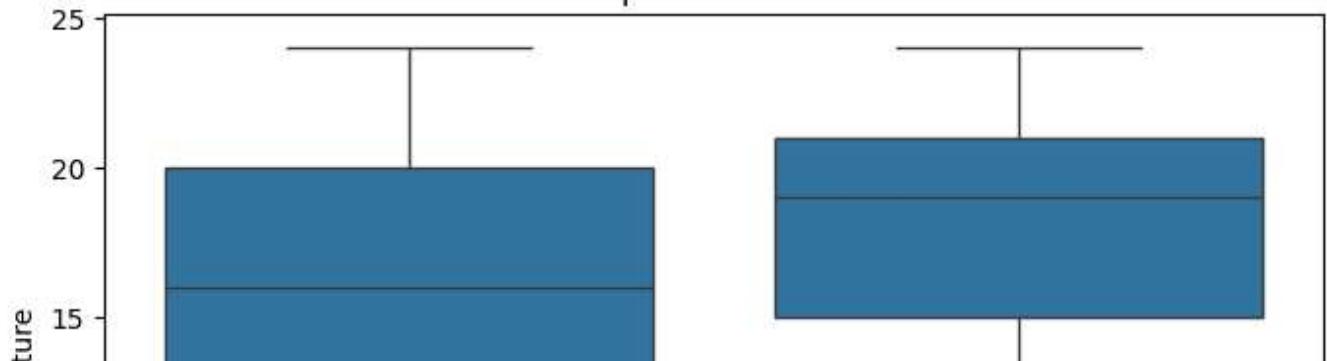


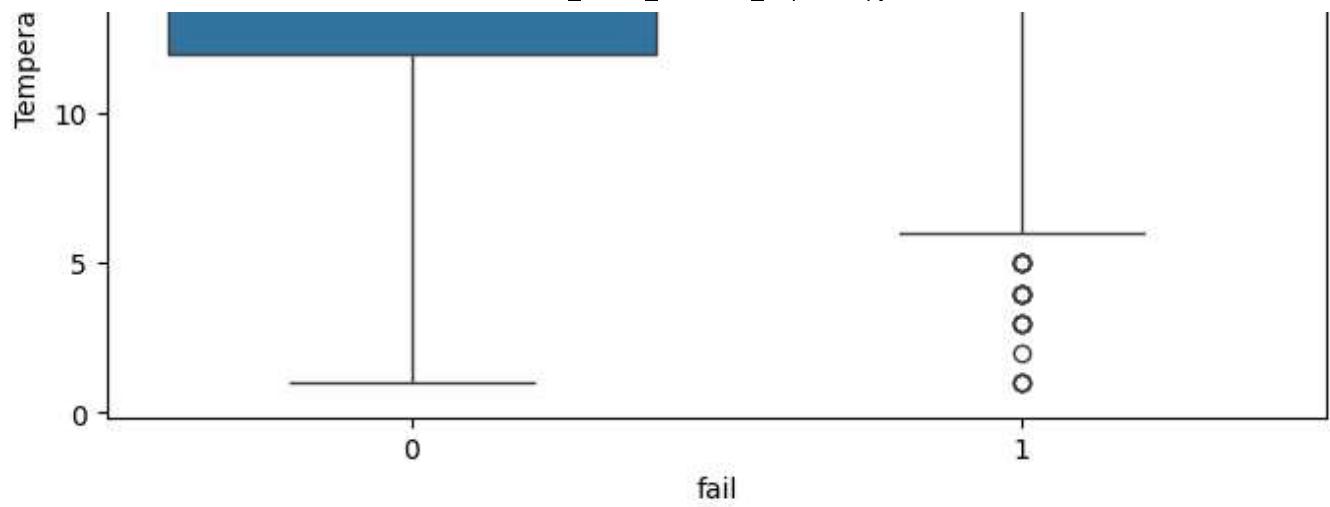


IP vs Fail



Temperature vs Fail



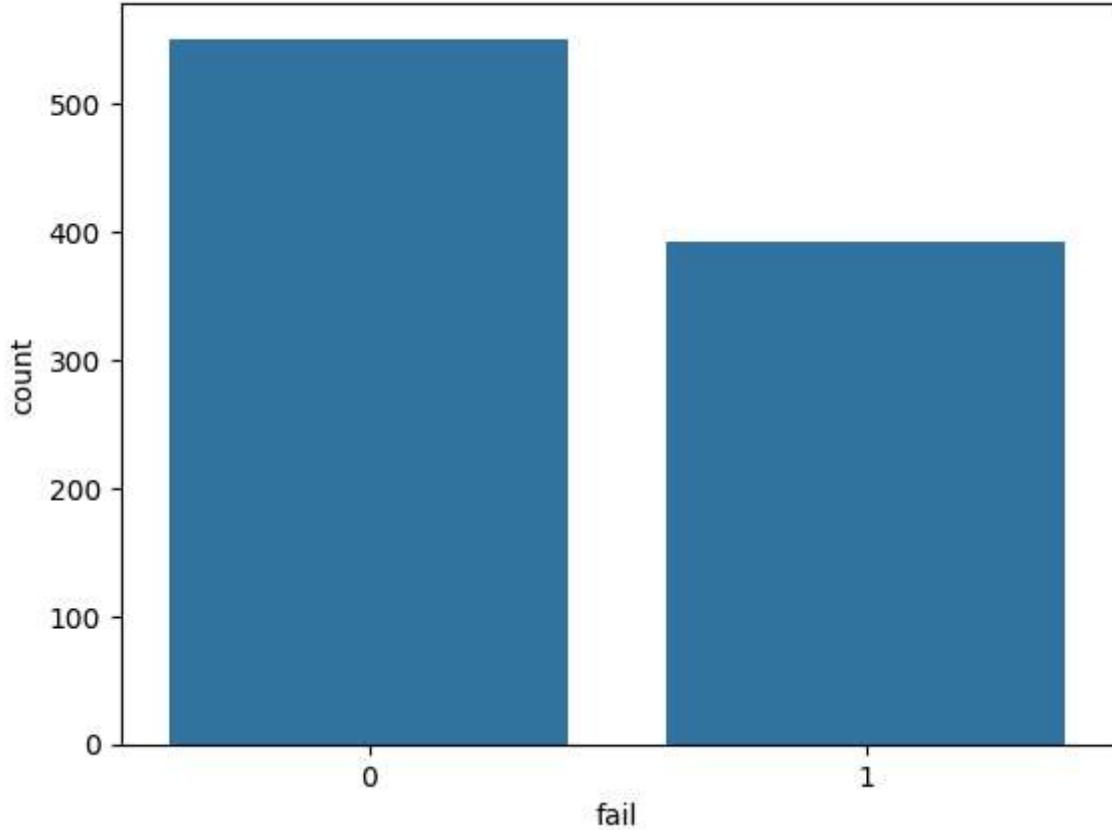


▼ Count plot of the target variable

```
# Count plot of the target variable
sns.countplot(x='fail', data=df)
plt.title('Failure Distribution')
plt.show()
```

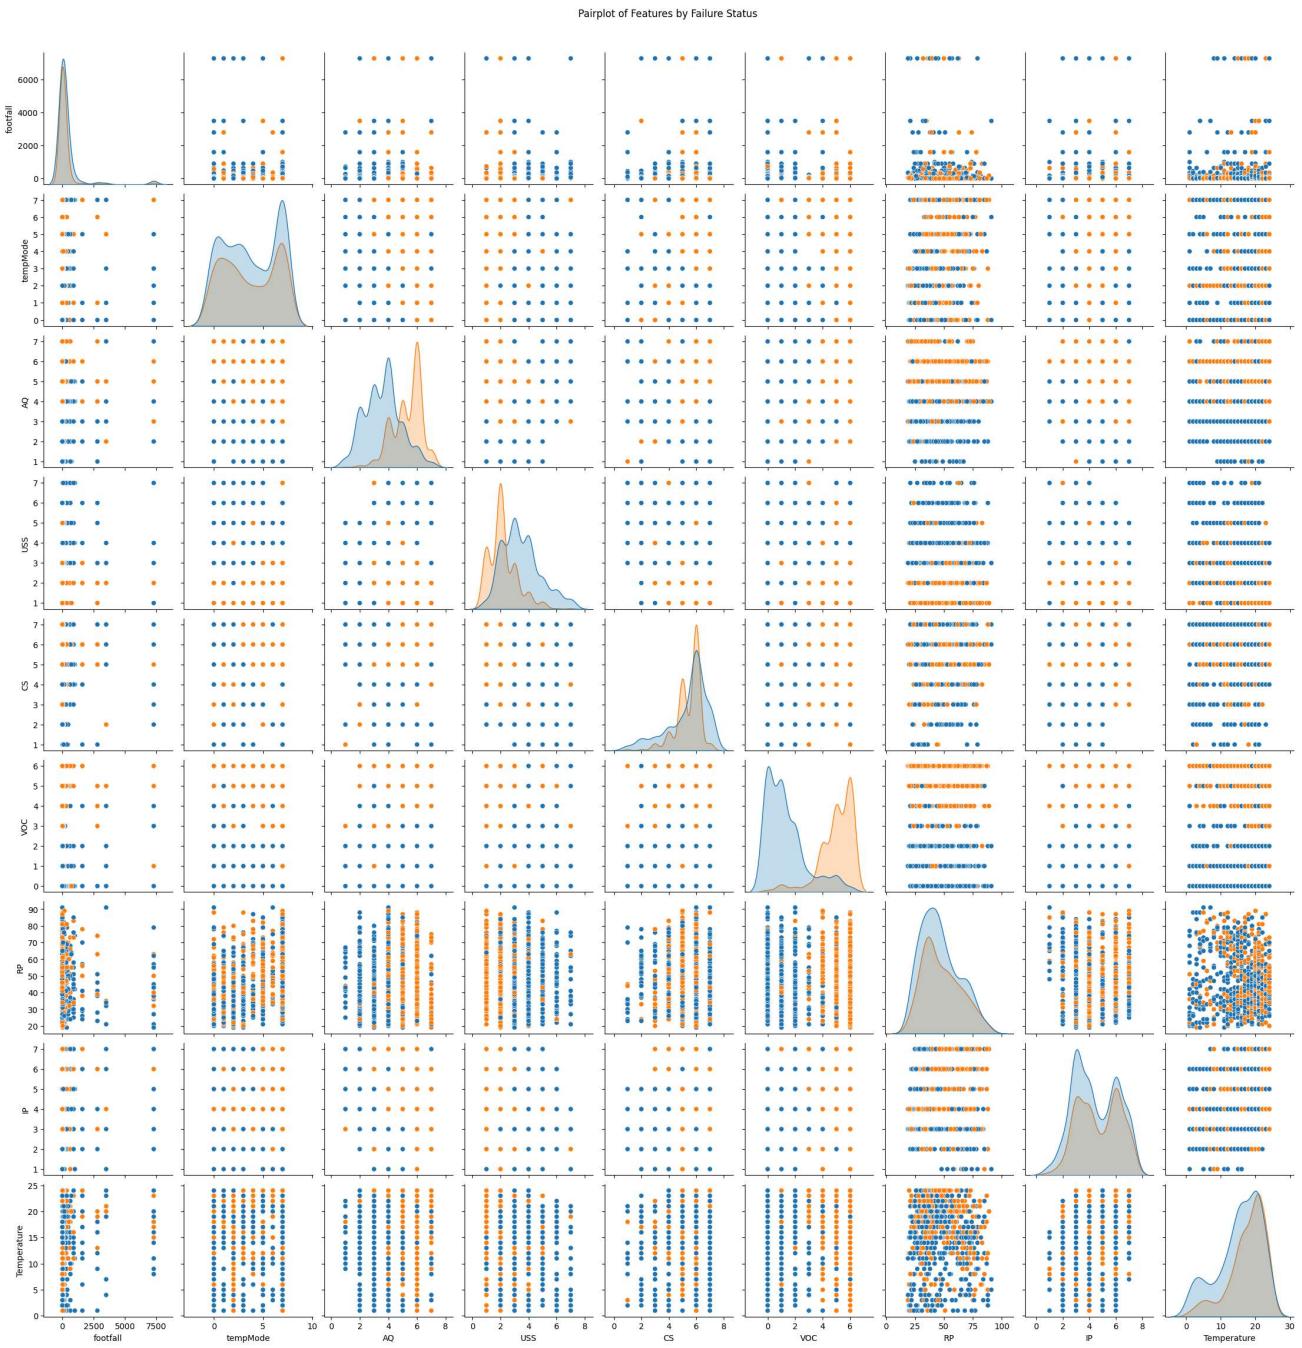


Failure Distribution



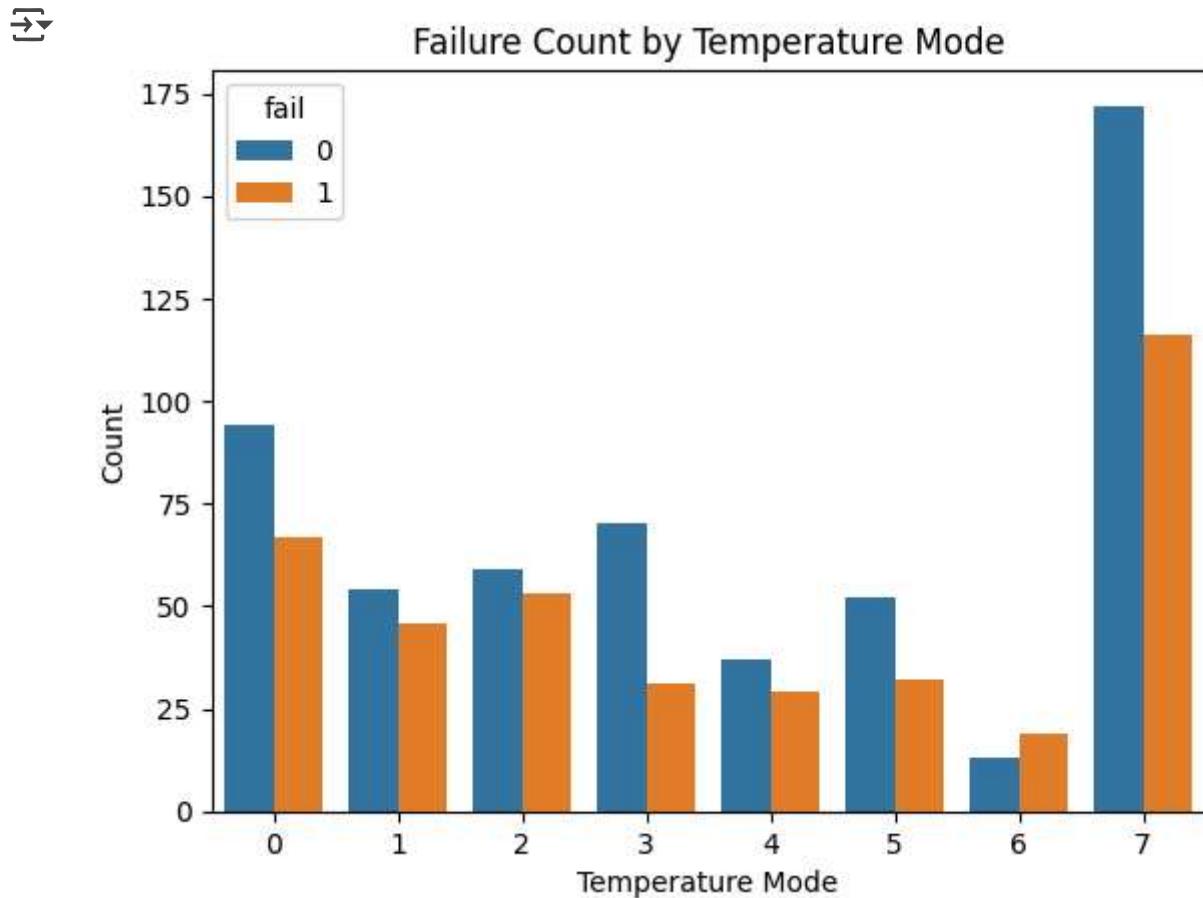
▼ Pairplot of numerical features colored by failure

```
# Pairplot of numerical features colored by failure
sns.pairplot(df.select_dtypes(include=['int64', 'float64']), hue='fail')
plt.suptitle('Pairplot of Features by Failure Status', y=1.02)
plt.show()
```



▼ Bar plot of 'tempMode' vs 'fail'

```
# Bar plot of 'tempMode' vs 'fail'  
sns.countplot(x='tempMode', hue='fail', data=df)  
plt.title('Failure Count by Temperature Mode')  
plt.xlabel('Temperature Mode')  
plt.ylabel('Count')  
plt.show()
```

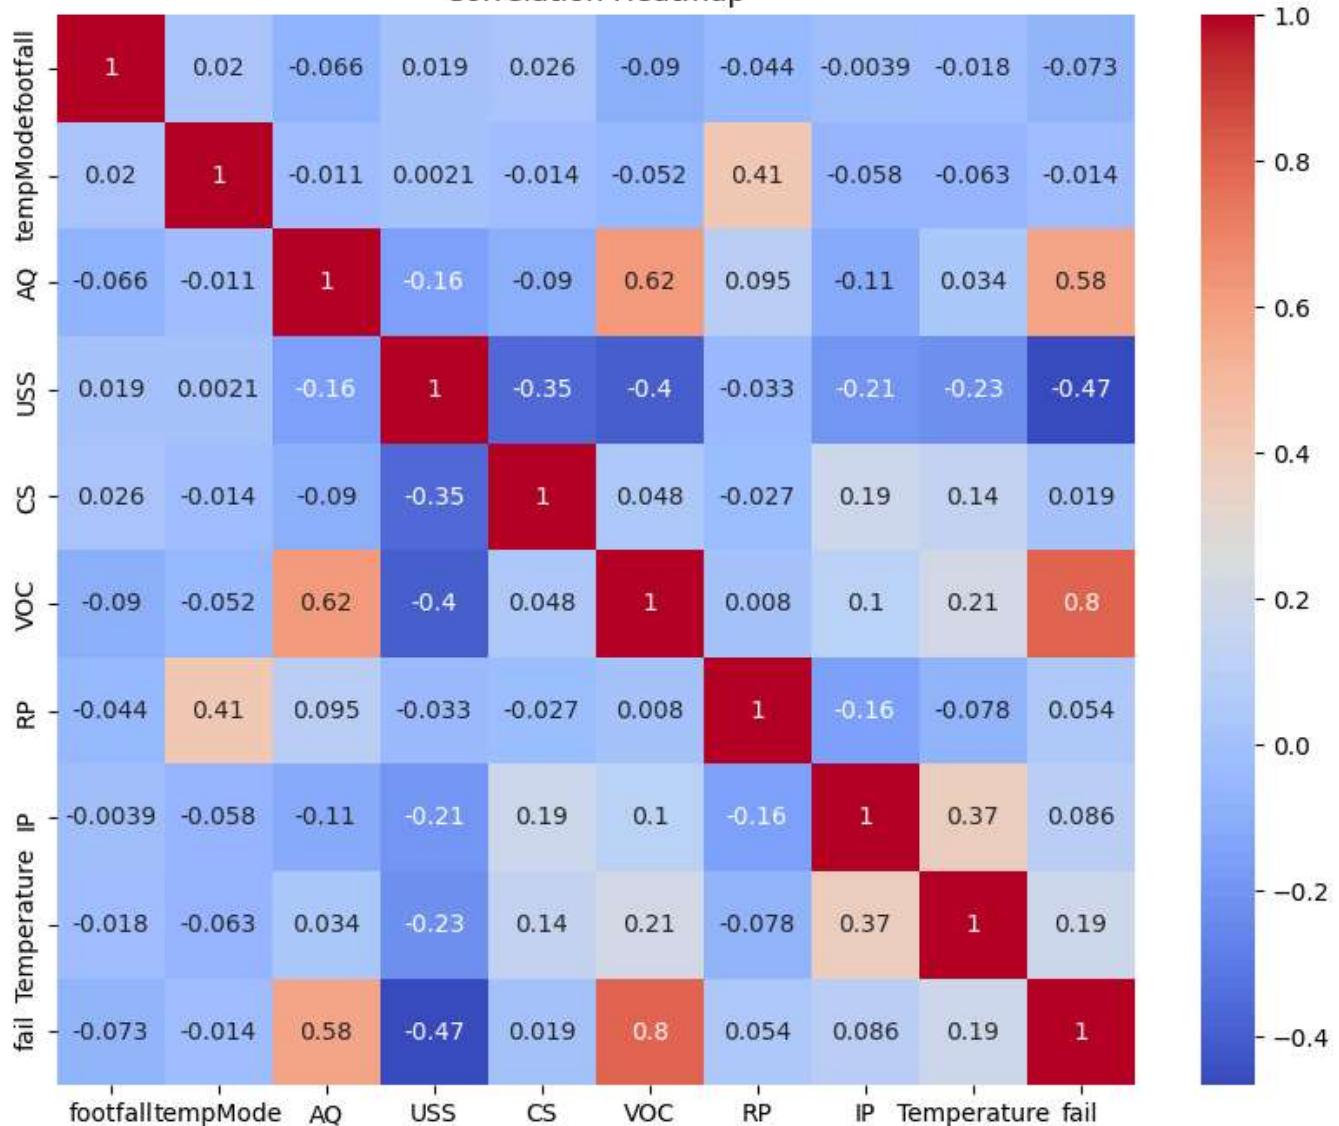


▼ Correlation Heatmap

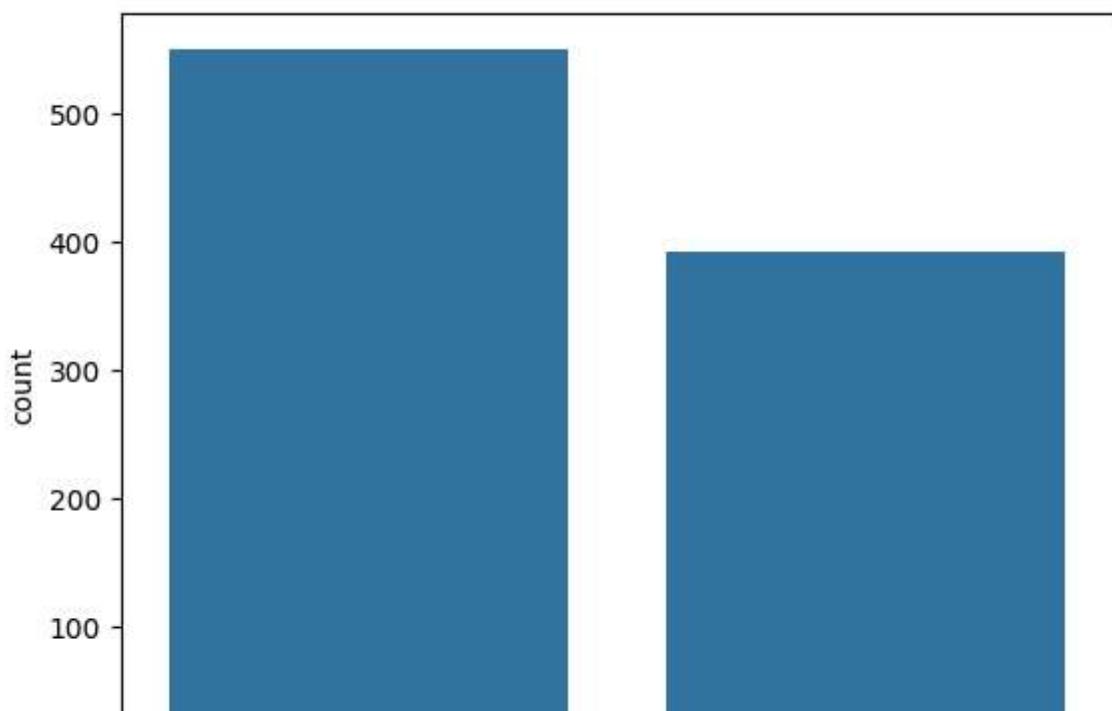
```
# Correlation Heatmap  
plt.figure(figsize=(10, 8))  
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')  
plt.title('Correlation Heatmap')  
plt.show()  
  
sns.countplot(x='fail', data=df)  
plt.title('Failure Distribution')  
plt.show()
```



Correlation Heatmap



Failure Distribution





❖ Feature Engineering & Scaling

❖ What is Feature Engineering?

Feature Engineering is the process of preparing and improving the dataset's input variables (**features**) so the machine learning model can understand and learn from them effectively.

In this step, we:

- **Separate features and target**
 - $X \rightarrow$ All sensor readings (input data)
 - $y \rightarrow$ Target column (`fail`) that indicates if the machine failed (1) or not (0)
- **Scale features** so they are on the same numerical range.

❖ Why Scaling?

Different features can have very different ranges. For example:

- **RPM** could be in thousands (e.g., 1500)
- **VOC levels** might be small decimal numbers (e.g., 0.05)

Without scaling, models like **Logistic Regression** and **XGBoost** might give more importance to features with larger values, even if they are not more important.

❖ How StandardScaler Works

StandardScaler transforms each numeric feature so it has:

- **Mean = 0**
- **Standard Deviation = 1**

Formula: $[z = \frac{x - \text{mean}}{\text{std}}]$

❖ Steps in Code

1. `X = df.drop('fail', axis=1)` → Drop the target column to create features dataset.
2. `y = df['fail']` → Create target vector for prediction.
3. Create a scaler: `scaler = StandardScaler()`
4. Apply scaling: `X_scaled = scaler.fit_transform(X)`

After scaling, all features are on a similar range, making the model training more stable and faster.

```
# Feature Engineering & Scaling
X = df.drop('fail', axis=1)
y = df['fail']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

▼ Model Training

▼ ↘ Purpose

In this step, we train our machine learning model to learn patterns from the scaled training data (`X_train`) and predict the target (`y_train`).

We are using: **Random Forest Classifier** – an ensemble learning method that:

- Creates multiple decision trees on random subsets of the data.
- Combines their predictions for a more accurate and stable result.
- Handles non-linear relationships and feature interactions well.

📌 Process

1. Train-Test Split

- `train_test_split()` splits the data into:
 - **Training set (80%)** → Used to teach the model.
 - **Testing set (20%)** → Used to evaluate the model on unseen data.
- `random_state=42` ensures results are reproducible.

2. Model Initialization & Training

- `RandomForestClassifier()` creates the model.
- `.fit(X_train, y_train)` trains it on the training data.

3. Prediction

- `.predict(X_test)` generates predictions for the unseen test set.

```
from sklearn.metrics import accuracy_score

# Split data
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42
)

print("Training set shape:", X_train.shape)
print("Testing set shape:", X_test.shape)

# Initialize and train model
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Quick accuracy check
accuracy = accuracy_score(y_test, y_pred)
print(f"✓ Model trained successfully! Accuracy on test set: {accuracy:.2f}")
```

→ Training set shape: (755, 9)
Testing set shape: (189, 9)
✓ Model trained successfully! Accuracy on test set: 0.88

▼ Logistic Regression model training

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

# Initialize Logistic Regression with useful parameters
model = LogisticRegression(
    random_state=42,           # Reproducibility
    max_iter=1000,             # Ensure convergence
    solver='liblinear',        # Works well for small to medium datasets
    penalty='l2',               # Regularization to avoid overfitting
    C=1.0                      # Regularization strength (lower = stronger regularization)
)

# Train the model
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Quick accuracy check
```

```
accuracy = accuracy_score(y_test, y_pred)
print("✅ Logistic Regression model training complete!")
print(f"📊 Accuracy on test set: {accuracy:.2f}")
```

→ ✅ Logistic Regression model training complete!
📊 Accuracy on test set: 0.87

▼ Make predictions on the test data

```
# Make predictions on the test data
y_pred = model.predict(X_test)

# Compare first 20 predictions with actual labels
comparison_df = pd.DataFrame({
    'Actual': y_test[:20].values,
    'Predicted': y_pred[:20]
})

print("\n🔍 First 20 Predictions vs Actual Values:")
print(comparison_df.to_string(index=False))

# Quick correctness check
correct = sum(comparison_df['Actual'] == comparison_df['Predicted'])
print(f"\n✅ Correct predictions out of first 20: {correct}/20")
```

→ 🔎 First 20 Predictions vs Actual Values:

Actual	Predicted
1	1
0	0
1	1
1	1
0	0
0	0
0	0
0	0
1	1
0	0
1	1
0	0
1	1
0	0
1	1
0	0
0	0
0	0

✅ Correct predictions out of first 20: 20/20

❖ Model Evaluation

❖ Purpose

Once the model is trained and predictions are made, we need to measure **how well it performed**. Model evaluation tells us whether the model is making accurate predictions and helps identify where it is going wrong.

📌 Key Metrics Used

1. Confusion Matrix

- A table showing how many predictions were:
 - **True Positives (TP)** → Correctly predicted failures
 - **True Negatives (TN)** → Correctly predicted non-failures
 - **False Positives (FP)** → Predicted failure, but it didn't fail (false alarm)
 - **False Negatives (FN)** → Predicted no failure, but it failed (missed failure)
- Helps us see if the model is more prone to false alarms or missed detections.

2. Classification Report

- **Precision** → Of the failures predicted, how many were correct?
Formula: $TP / (TP + FP)$
- **Recall (Sensitivity)** → Of the actual failures, how many did we catch?
Formula: $TP / (TP + FN)$
- **F1 Score** → Balance between Precision and Recall.
- **Accuracy** → Overall % of correct predictions.